



Expert-level validation of AI-generated medical text

Asad Aali
Talk, Daneshjou Lab
July 22, 2025

About me

- Research Scientist at Stanford University
 - Lab: Machine Intelligence for Medical Imaging (MIMI)
 - Advisor: Akshay Chaudhari
- Passionate about developing machine learning methods for healthcare applications
- Research Interests:
 - Machine Learning
 - Foundation Models
 - Healthcare



Motivation

- Language models (LMs) are increasingly leveraged for documentation in medical settings
- However, they frequently generate **plausible text** that contains **subtle errors**
- **GOAL:** Develop validators that perform **scalable** and **reliable risk assessment**
 - Validate whether **AI-generated medical text** is safe for **patient-facing use cases**, specifically in **radiology report** and **discharge summary simplification (binary assessment)**.
- **Medical text validation** determines whether the output is **factually consistent** with the input
Radiology Impression (*Input*) -> Patient-Friendly Impression (*Output*)

Summary:

- Development of **scalable, automated tools** for evaluating LM-generated medical texts.
- Ensuring the tools integrate **clinical expertise** and **stakeholder feedback**.
- Expanding and optimizing LMs for assessing outputs in **radiology reports** and **discharge summaries**.
- Creating **synthetic datasets** and robust evaluation frameworks to improve LM reliability.

Related Works

- The "LM-as-judge" paradigm, has shown promise in automating the evaluation of LM-generated text.
- Most general-purpose LM-as-judge methods treat LMs as **static evaluators**
 - often **lacking the granularity** required to assess nuanced, high-stakes medical text
- Recent **medical-specific** LM-as-judge approaches suffer from certain **practical limitations**
 - assume resources like **expert-labeled training data**
 - assume availability of **reference outputs** or **retrieval-based evidence**
- Many approaches focus on **subdomains** such as radiology (e.g., chest X-rays)

Method	Focus	Medical domain	Train -able	Physician -free training	Reference -free evaluation	Retrieval -free evaluation	Multi -lingual evaluation
FActScore	General	✗	✓	✓	✓	✗	✗
AlignScore	General	✗	✓	✓	✓	✓	✗
FineRadScore	Radiology	✓	✗	✓	✗	✓	✗
ReXTrust	Radiology	✓	✓	✗	✗	✓	✗
GREEN	Radiology	✓	✓	✓	✗	✓	✗
VeriFact	BHC	✓	✗	✓	✓	✗	✗
DocLens	Medical	✓	✗	✓	✗	✓	✗
MedHAL	Medical	✓	✓	✗	✗	✓	✗
MedVAL	Medical	✓	✓	✓	✓	✓	✓

Introduction

- MedVAL = Medical Text Validator (not Evaluator)
- A **self-supervised** framework that leverages **synthetic data** to train LMs for robust medical text validation
 - Involves curating **high-quality** synthetic training examples by leveraging the **agreement between a generator and a validator LM as a proxy for physician judgment**
- MedVAL assesses whether an **output is factually consistent with the input** by assigning one of **four risk levels**, while flagging "**unsafe for deployment**" outputs at **near physician-level reliability**
- MedVAL extends prior methods by providing a framework that:
 - trains scalable evaluators **without physician-in-loop supervision**,
 - assesses medical text in the **absence of reference outputs or retrieval**,
 - supports **multilingual evaluation**,
 - offers **interpretable, expert-aligned assessments**.

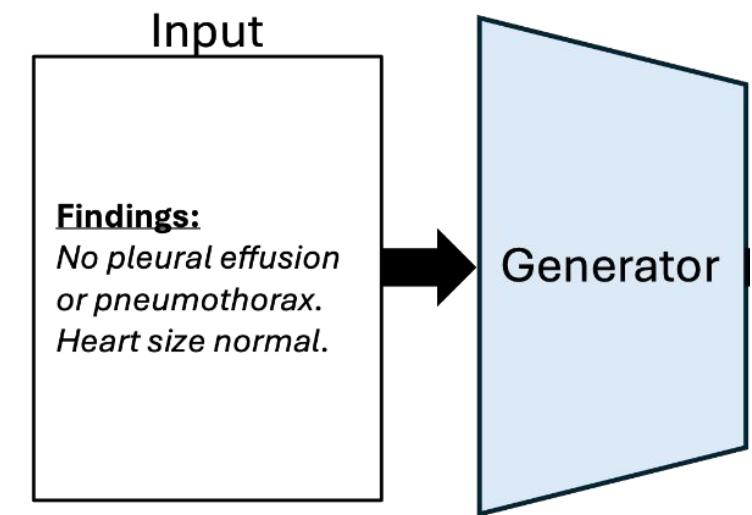
MedVAL Workflow

Input

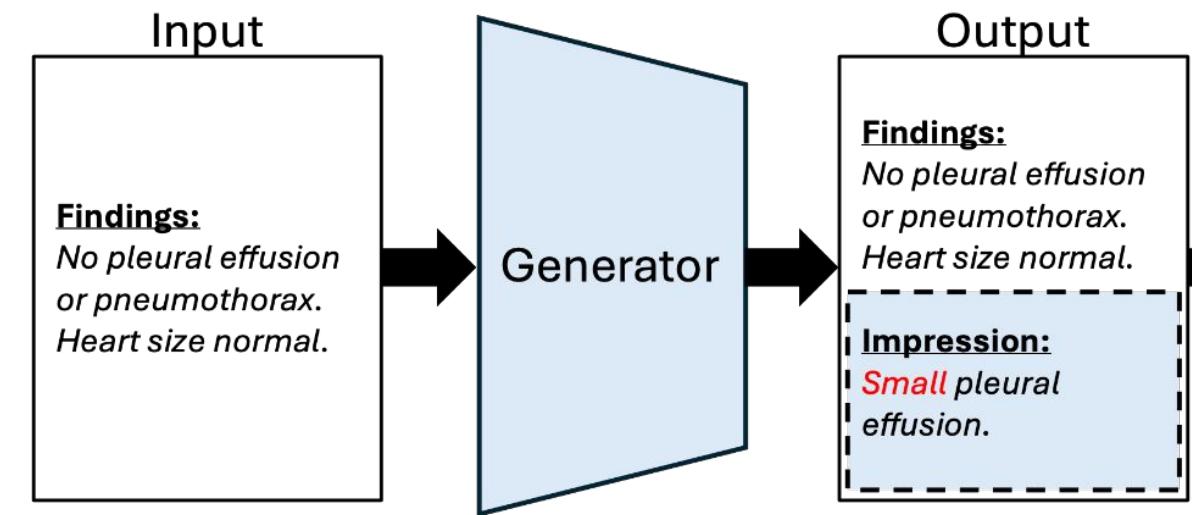
Findings:

*No pleural effusion
or pneumothorax.
Heart size normal.*

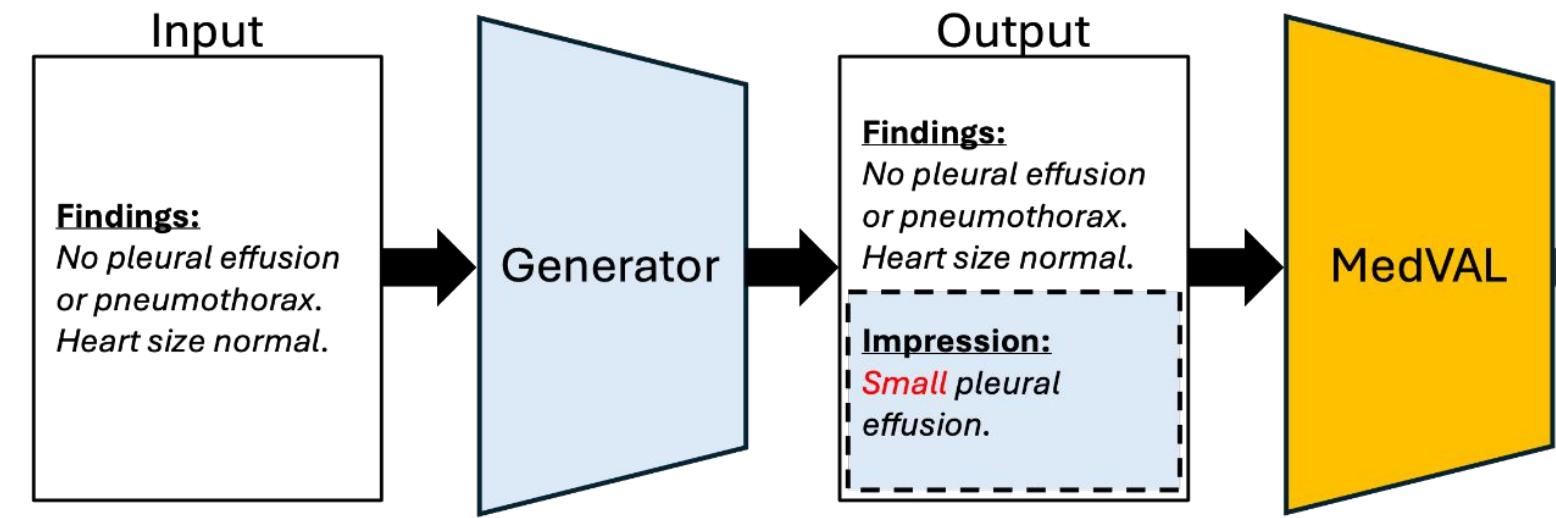
MedVAL Workflow



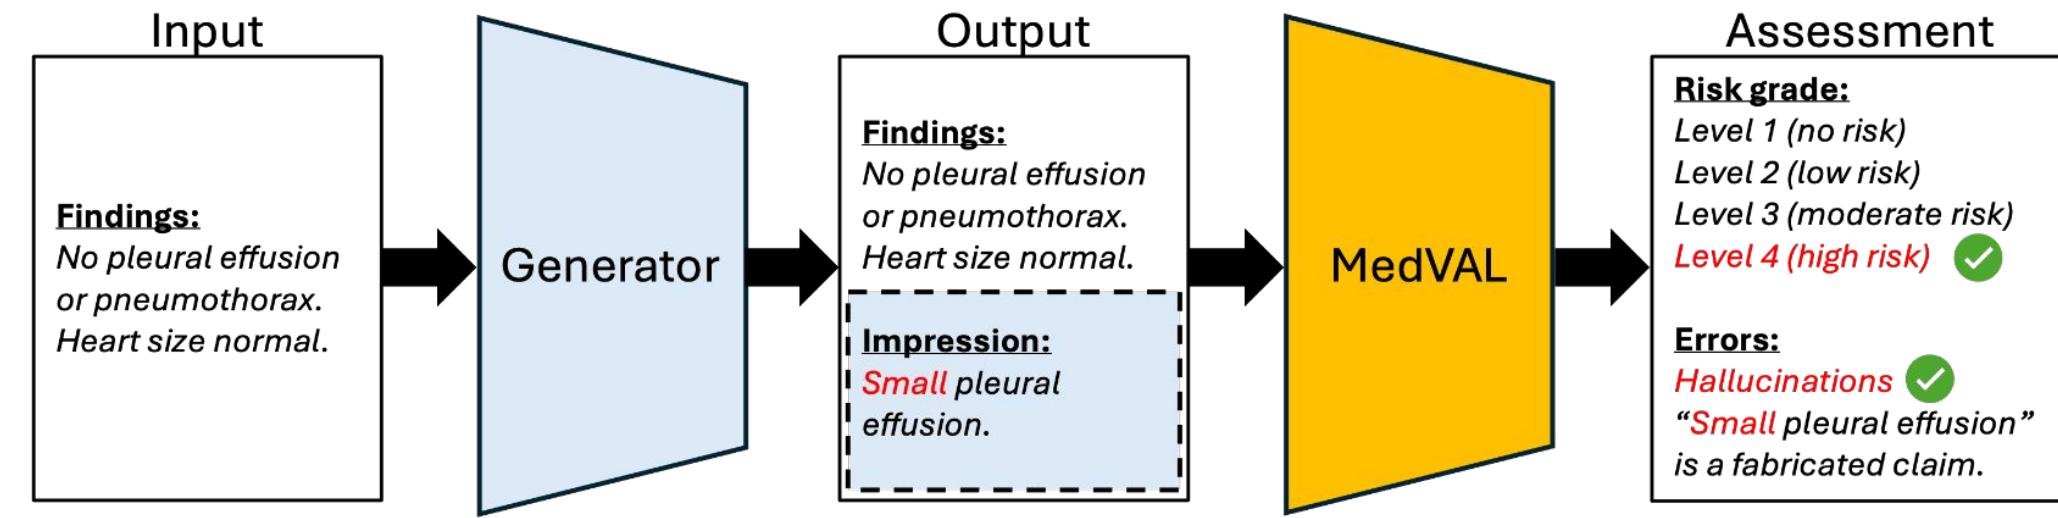
MedVAL Workflow



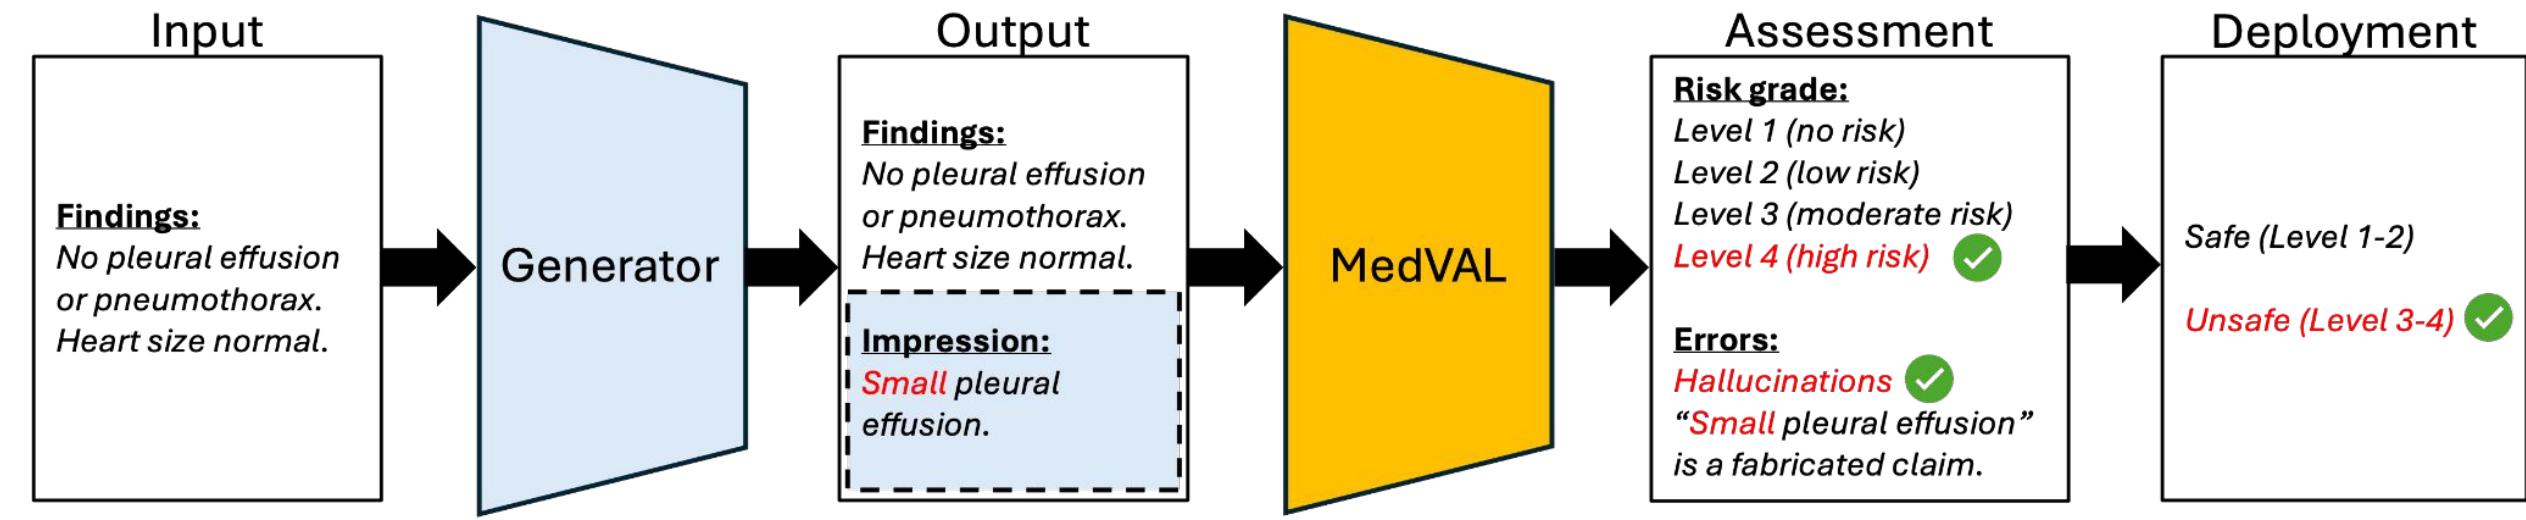
MedVAL Workflow



MedVAL Workflow



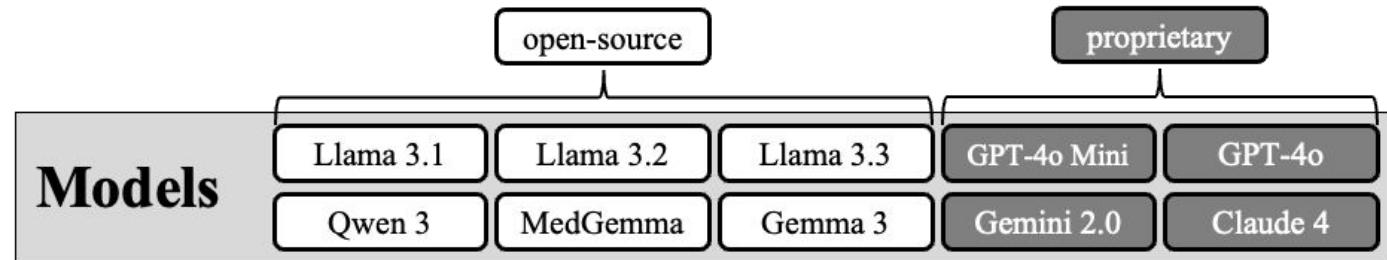
MedVAL Workflow



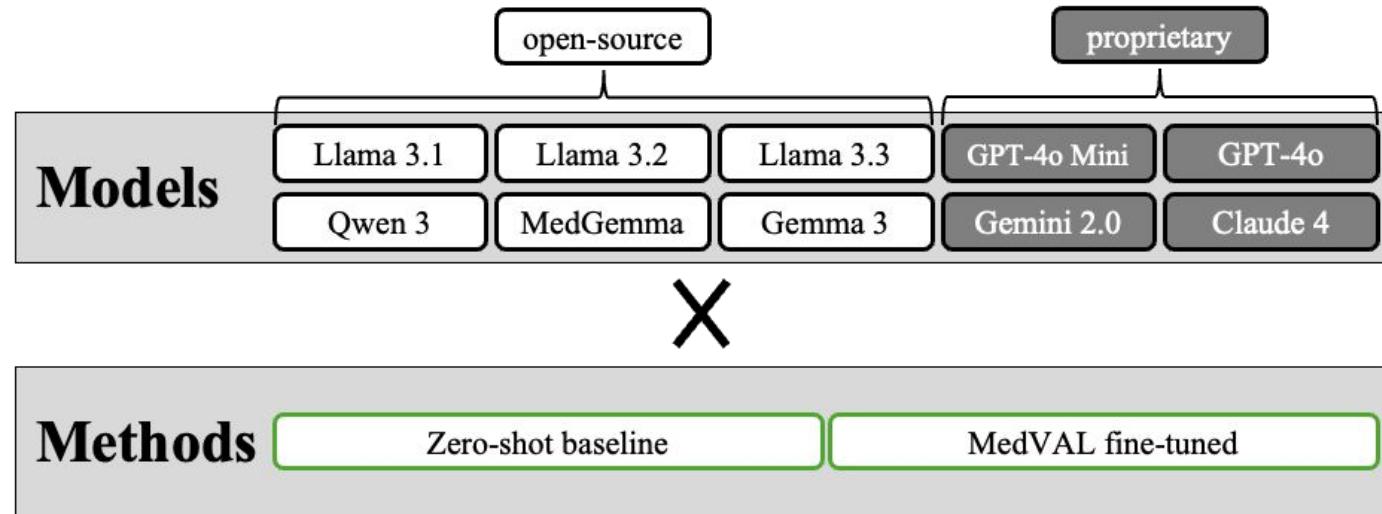
Contributions

1. MedVAL self-supervised training:
 - A general-purpose, self-supervised framework for training LMs to **validate factual consistency**
2. MedVAL-Bench dataset (Medical Text Validation Benchmark):
 - **GOAL:** Enabling rigorous evaluation of alignment b/w LMs and physicians in medical text validation
 - A dataset containing **840 physician-labeled evaluations** of AI-generated medical text
 - Performed by **12 physicians**
 - Spanning **6 diverse** medical text generation tasks
3. MedVAL performance benchmark:
 - MedVAL fine-tuning improves the validation **capabilities of all underlying LMs** across **in-distribution** and **out-of-distribution** tasks
 - MedVAL yields **significant gains** ($p < 0.001$): average baseline F1 scores for:
 - **four-class risk-level classification** improve from **36.7% to 51.0%**
 - **binary safe/unsafe classification** improve from **66.2% to 82.8%**

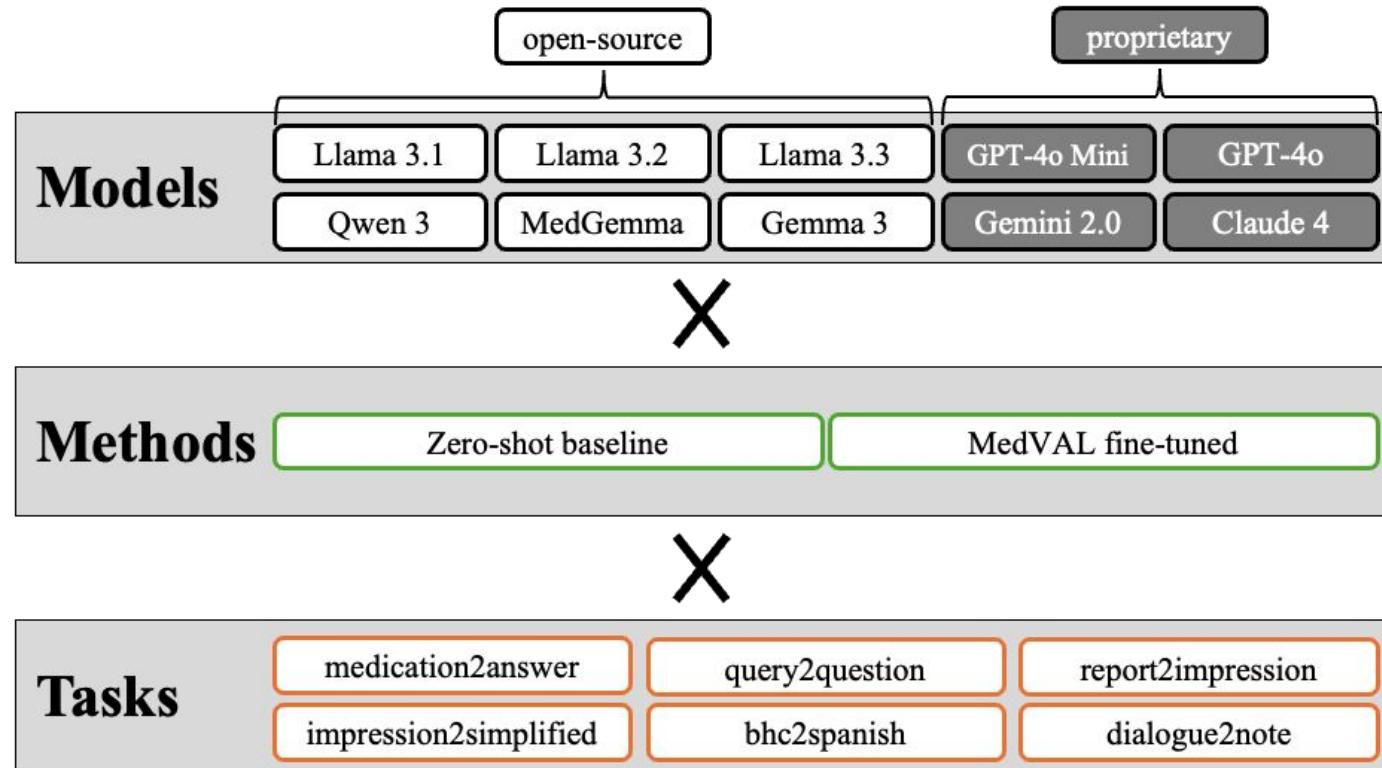
Schematic



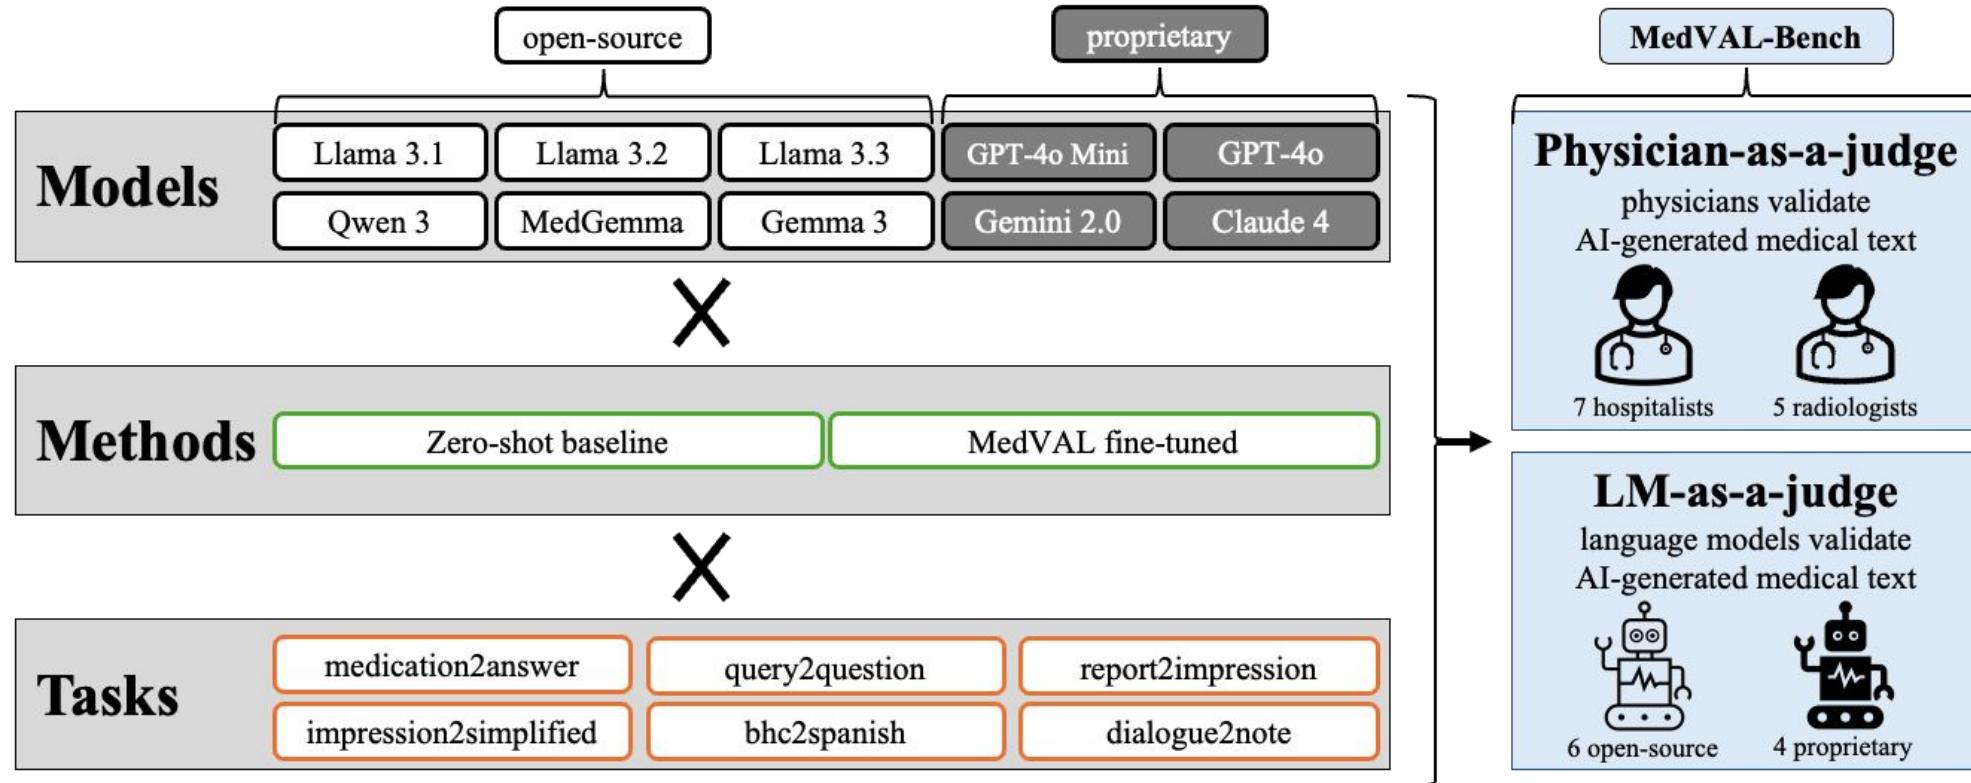
Schematic



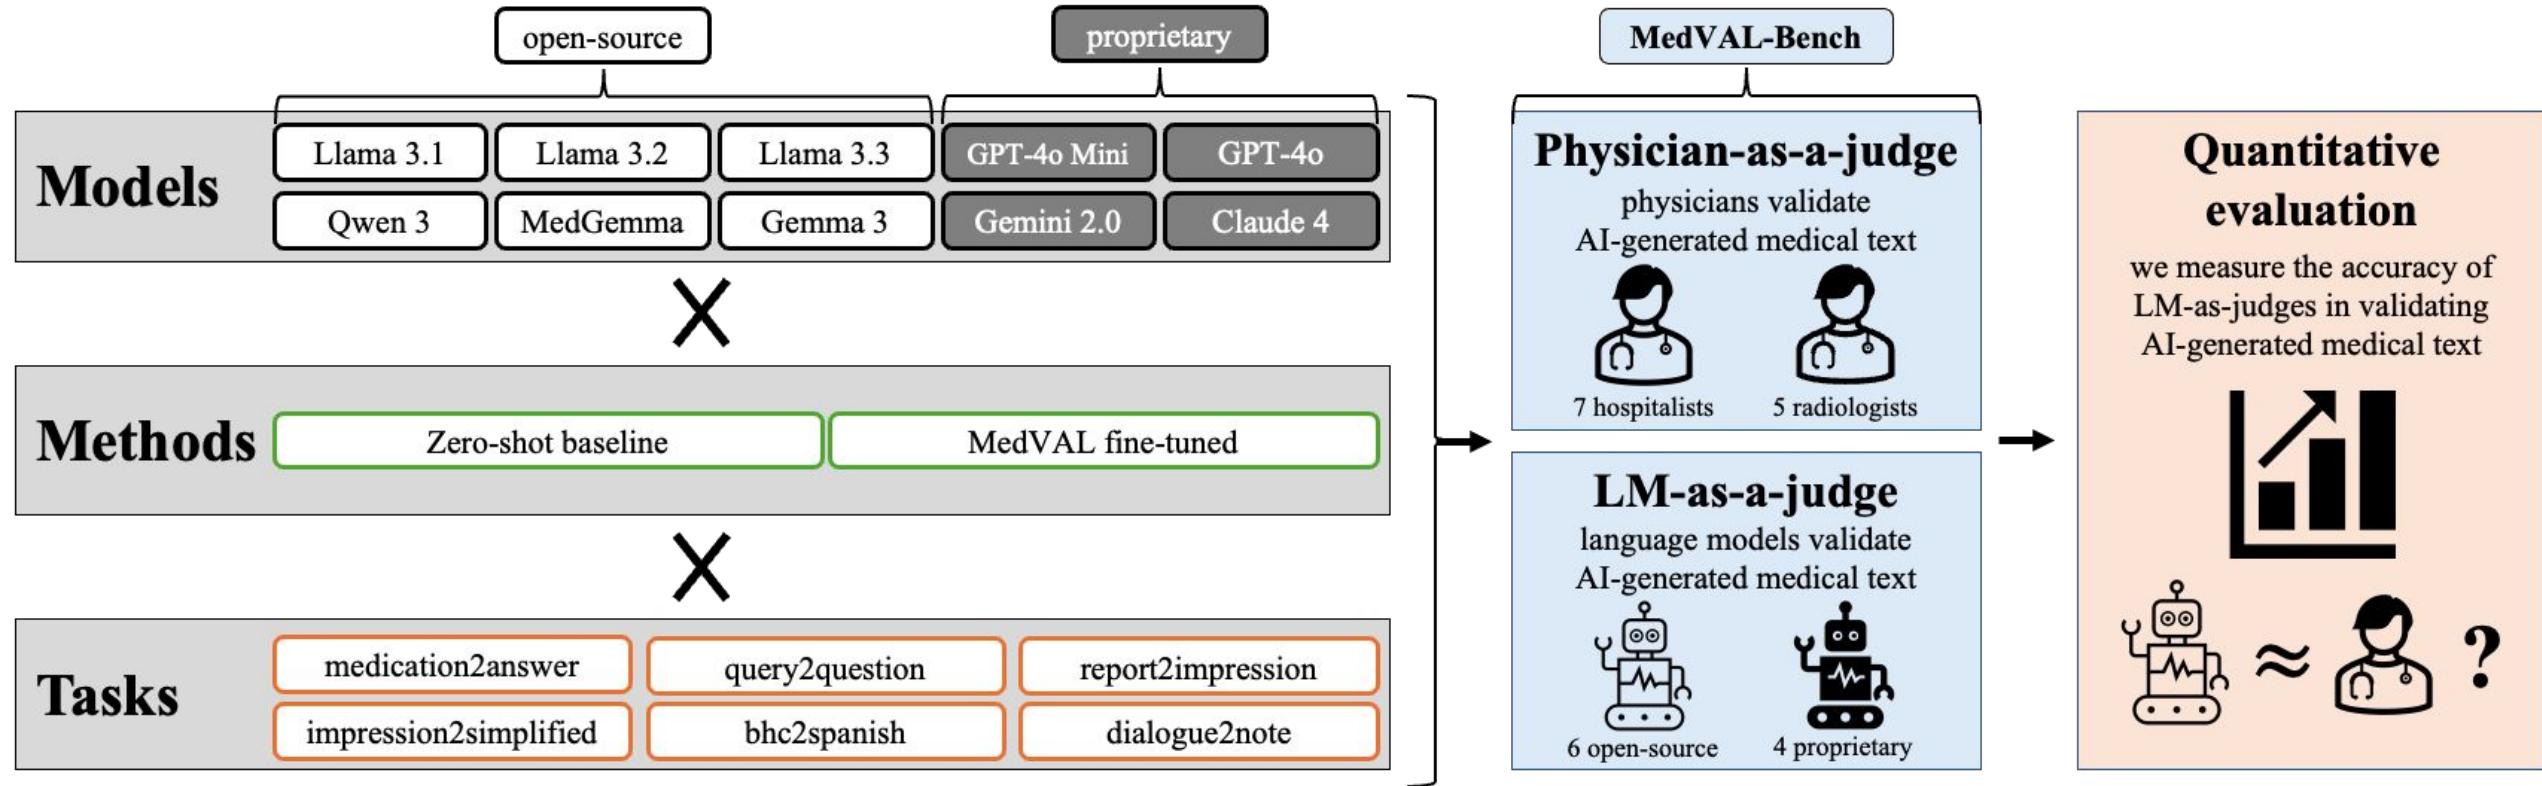
Schematic



Schematic



Schematic



MedVAL Training

Stage 1: Synthetic data generation

Input x

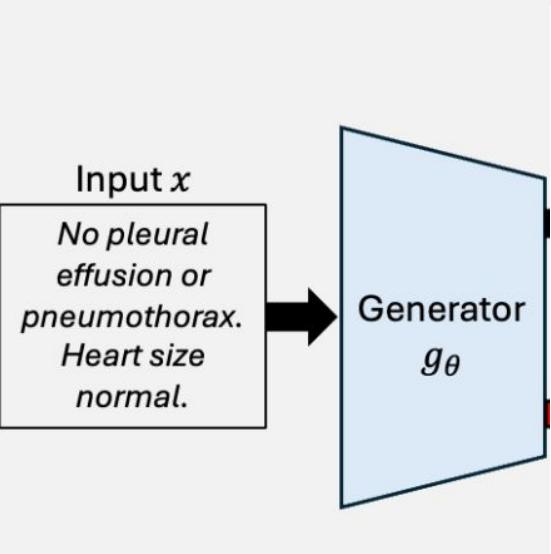
*No pleural effusion or pneumothorax.
Heart size normal.*

Stage 2: Data filtering

MedVAL Training

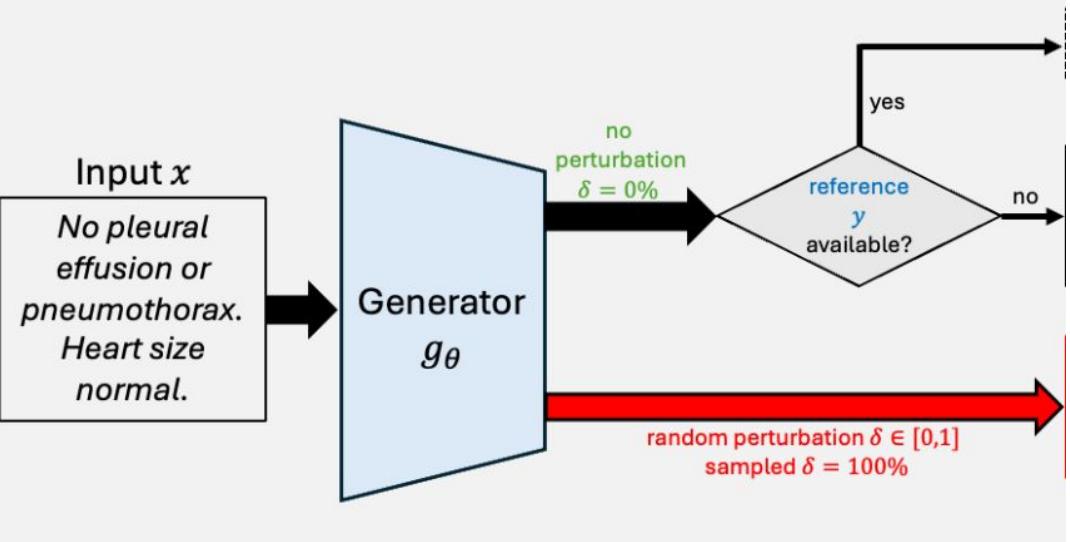
Stage 1: Synthetic data generation

Stage 2: Data filtering



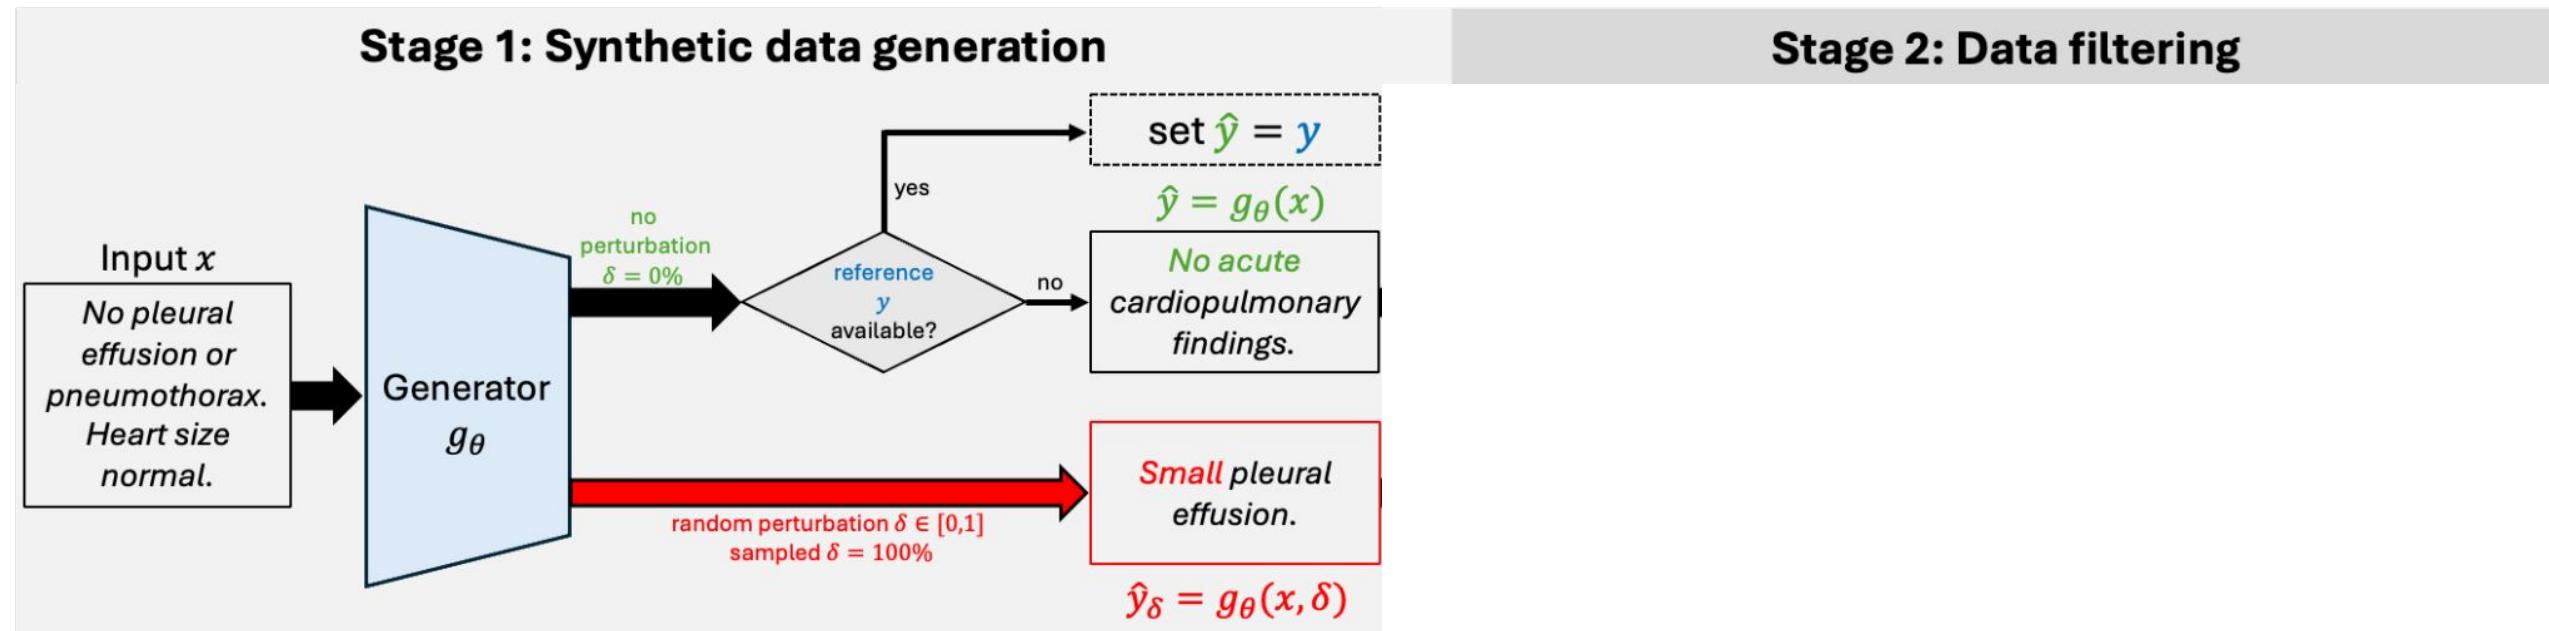
MedVAL Training

Stage 1: Synthetic data generation

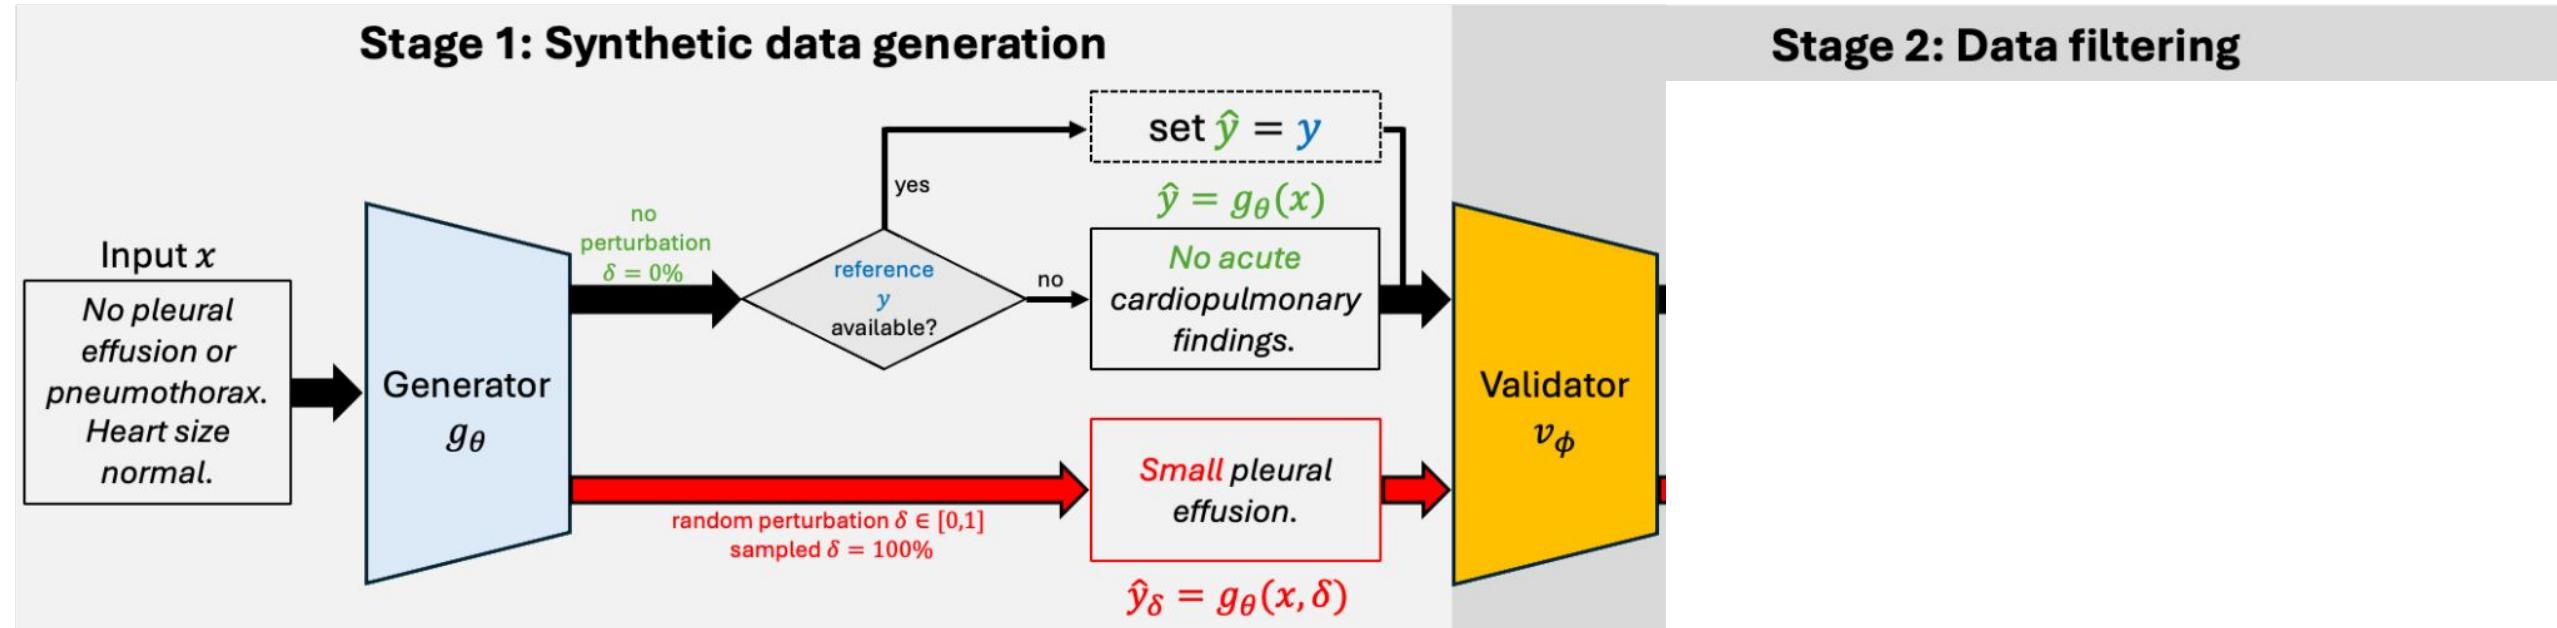


Stage 2: Data filtering

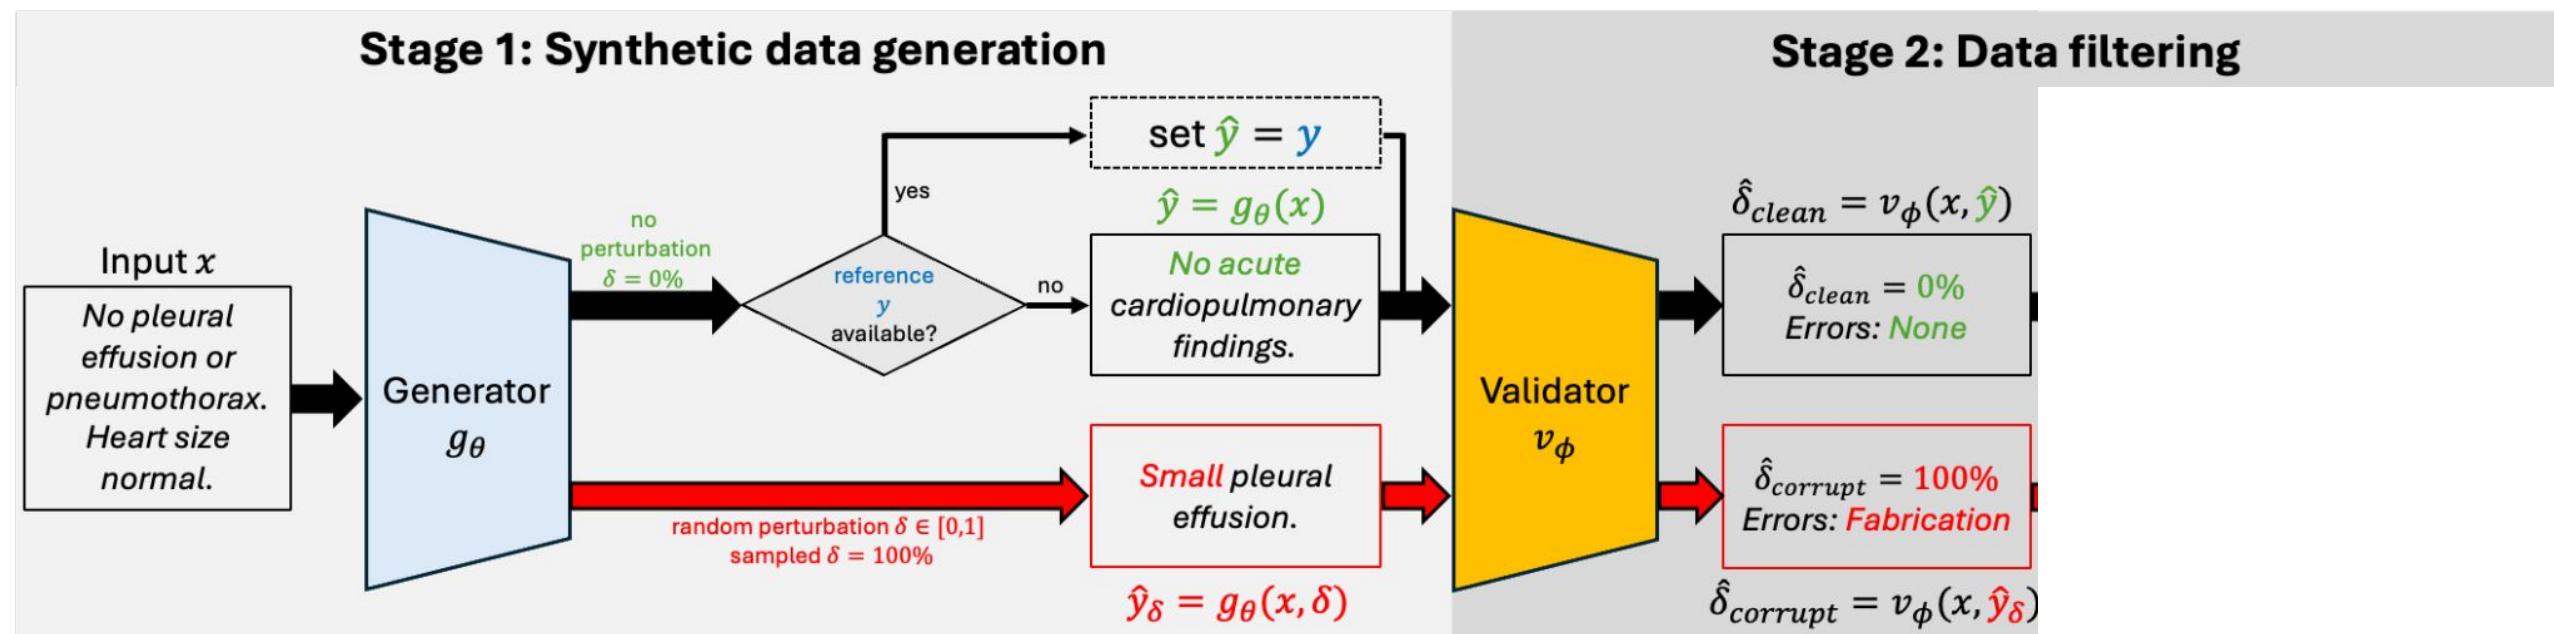
MedVAL Training



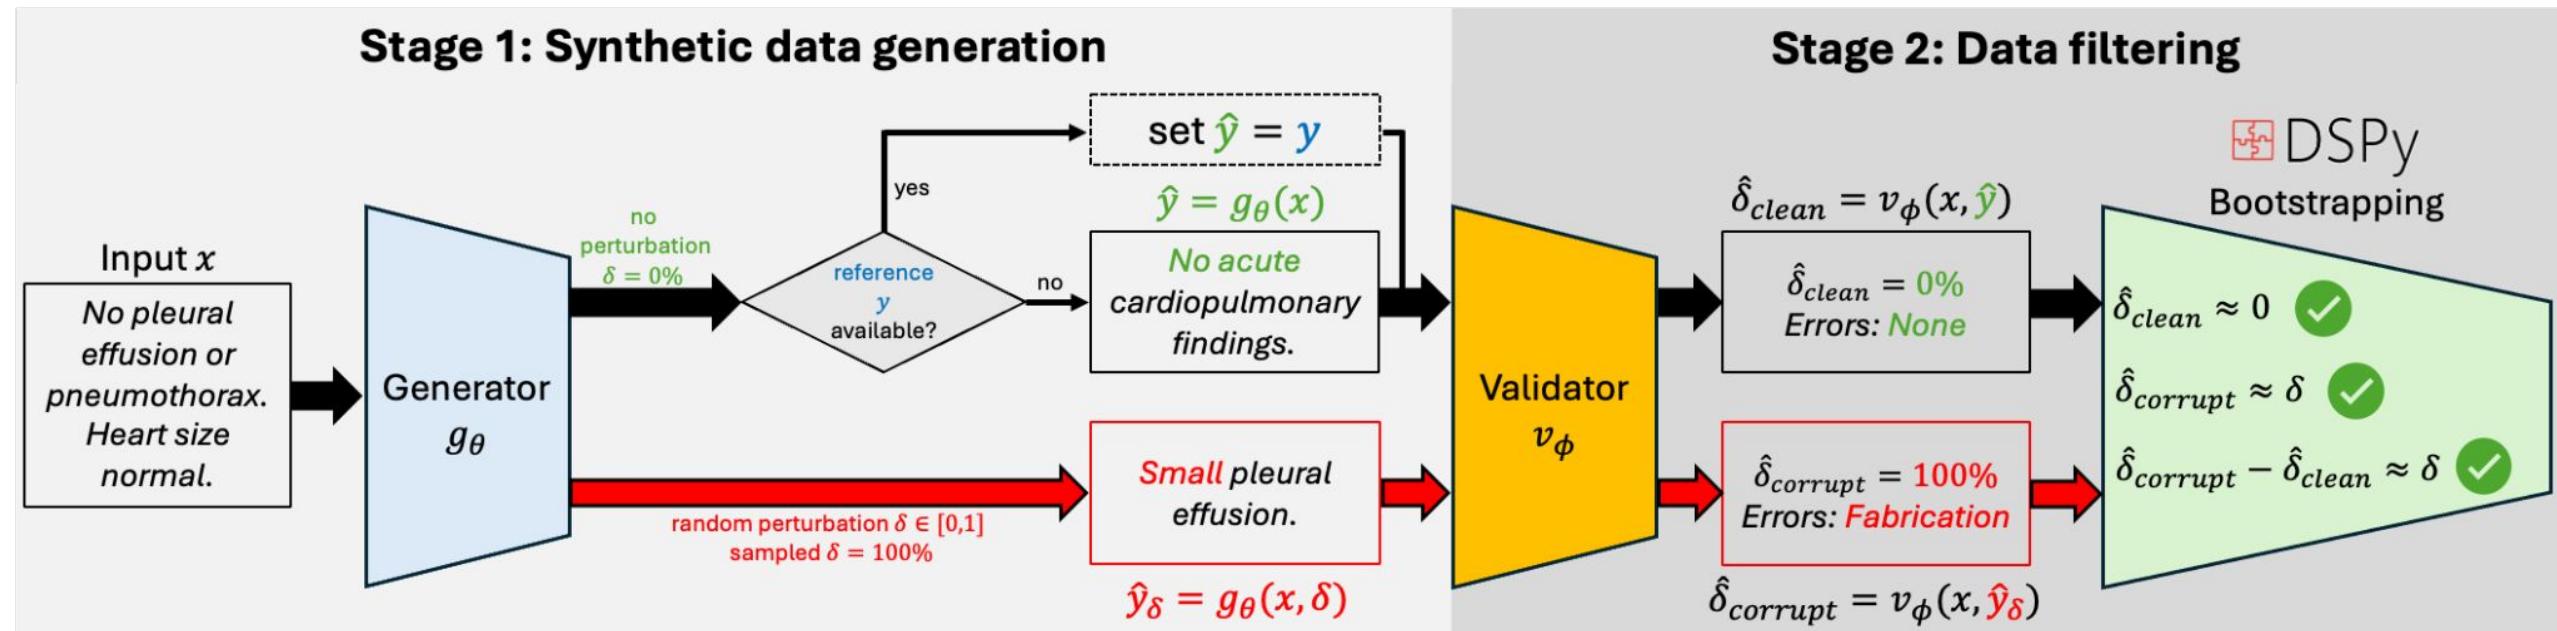
MedVAL Training



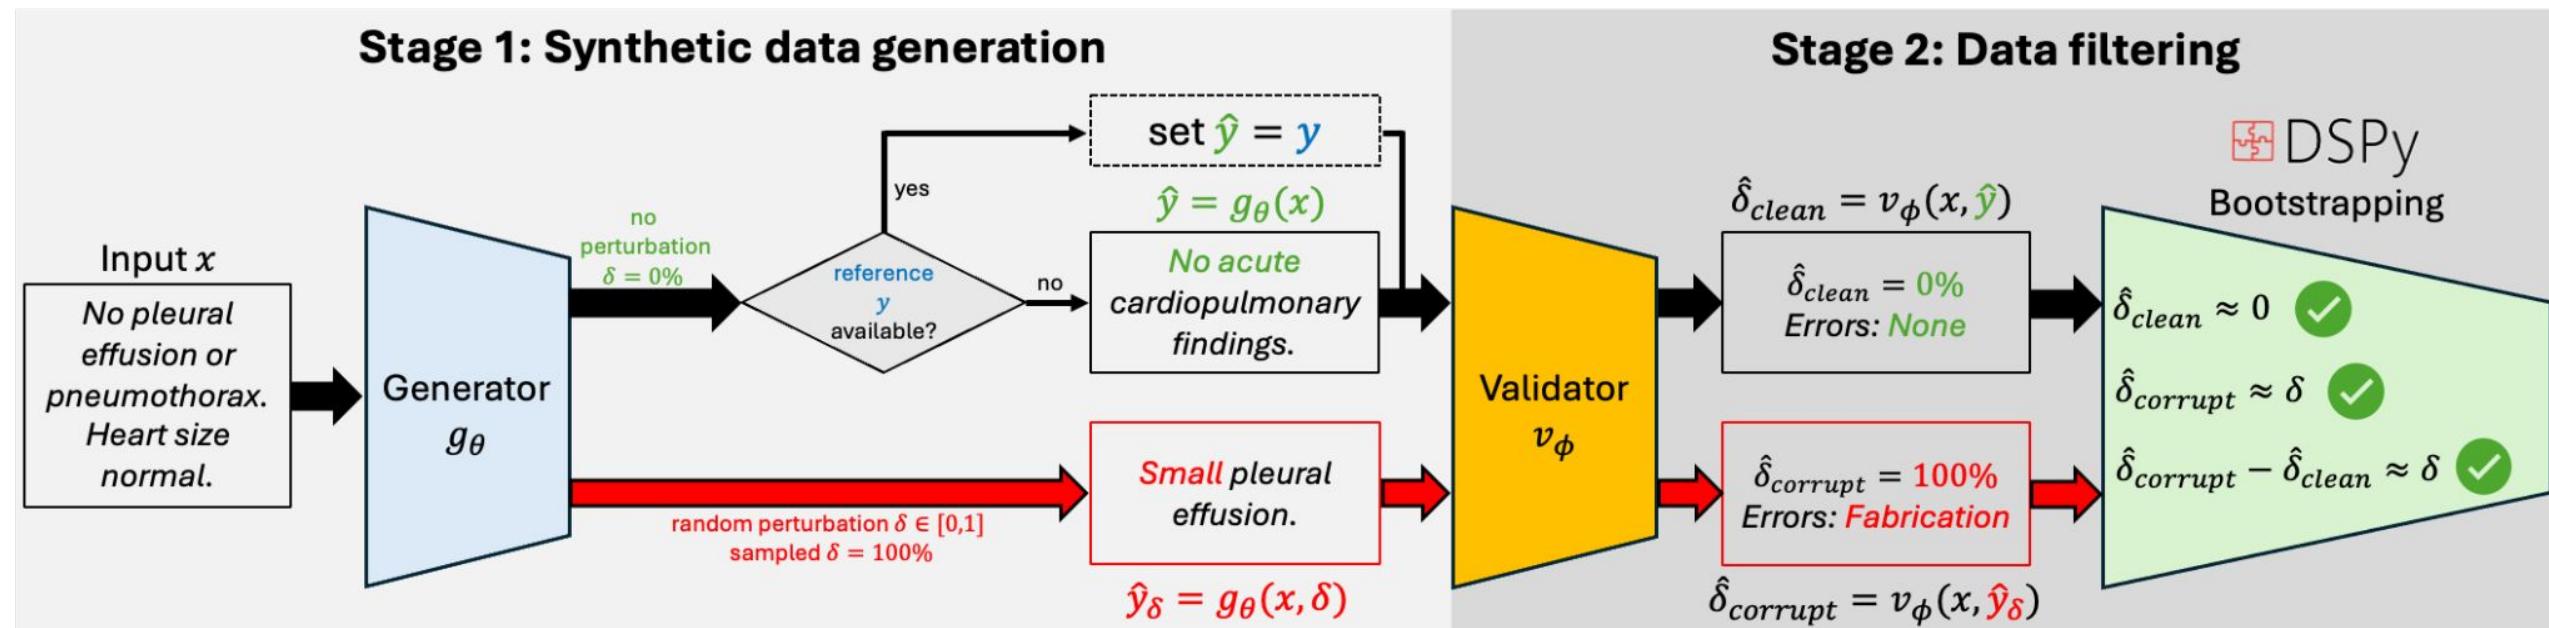
MedVAL Training



MedVAL Training



MedVAL Training



Perturbation	Category	Risk	Safety	Action
$\delta = 0\%$	Level 1	→ No Risk	Safe	Expert review not required.
$\delta = 33\%$	Level 2	→ Low Risk	Acceptable	Expert review optional.
$\delta = 67\%$	Level 3	→ Moderate Risk	Potentially unsafe	Expert review required.
$\delta = 100\%$	Level 4	→ High Risk	Unsafe	Expert rewrite required.

MedVAL Training - Algorithm

Algorithm 1 MedVAL self-supervised training

Require: Frozen generator g_θ , frozen validator v_ϕ , fine-tunable validator v_α , inputs $\mathcal{D} = \{x_i\}$, threshold τ

Ensure: Trained validator v_α^*

```

1: Initialize training dataset  $\mathcal{D}_{\text{train}} \leftarrow \emptyset$ 
2: for  $x \in \mathcal{D}$  do
3:    $\delta \leftarrow \text{RandomChoice}(\{\delta_1, \delta_2, \dots, \delta_L\} \mid \delta \in [0, 1])$ 
4:    $\hat{y} \leftarrow y$  if available, else  $g_\theta(x)$                                 ▷ Unperturbed output
5:    $\hat{y}_\delta \leftarrow g_\theta(x_\delta)$                                          ▷ Perturbed output
6:    $\hat{\delta}_{\text{clean}} \leftarrow v_\phi(x, \hat{y})$                                ▷ Factual degradation of  $\hat{y}$  in comparison to  $x$ 
7:    $\hat{\delta}_{\text{corrupt}} \leftarrow v_\phi(x, \hat{y}_\delta)$                            ▷ Factual degradation of  $\hat{y}_\delta$  in comparison to  $x$ 
8:   Compute  $\mathcal{M}_{\text{absolute}} \leftarrow \|\hat{\delta}_{\text{clean}}\|_2^2 + \|\hat{\delta}_{\text{corrupt}} - \delta\|_2^2$     ▷ Absolute consistency
9:   Compute  $\mathcal{M}_{\text{relative}} \leftarrow \|\hat{\delta}_{\text{corrupt}} - \hat{\delta}_{\text{clean}} - \delta\|_2^2$            ▷ Relative consistency
10:   $\mathcal{M}_{\text{MedVAL}} \leftarrow 1 - \frac{1}{6}(\mathcal{M}_{\text{absolute}} + \mathcal{M}_{\text{relative}})$           ▷ Generator-validator consistency score (0-1)
11:  if  $\mathcal{M}_{\text{MedVAL}} \geq \tau$  then
12:     $\mathcal{D}_{\text{train}} \leftarrow \mathcal{D}_{\text{train}} \cup \{x, \hat{y}, \hat{\delta}_{\text{clean}}\}$ 
13:     $\mathcal{D}_{\text{train}} \leftarrow \mathcal{D}_{\text{train}} \cup \{x, \hat{y}_\delta, \hat{\delta}_{\text{corrupt}}\}$ 
14:  end if
15: end for
16:  $v_\alpha^* = SFT(v_\alpha, \mathcal{D}_{\text{train}})$                                          ▷ Supervised fine-tuning
17: return  $v_\alpha^*$ 

```

Perturbation Strategies

Perturbation	Instructional prompt
$\delta = 0\%$	"The output should contain no clinically meaningful factual inconsistencies. Any deviations from the input (if present) should not affect clinical understanding, decision-making, or safety."
$\delta = 33\%$	"The output should contain subtle or ambiguous inconsistencies that are unlikely to influence clinical decisions or understanding. These inconsistencies should not introduce confusion or risk."
$\delta = 67\%$	"The output should contain inconsistencies that could plausibly affect clinical interpretation, documentation, or decision-making. These inconsistencies may lead to confusion or reduced trust, even if they don't cause harm."
$\delta = 100\%$	"The output should include one or more inconsistencies that could result in incorrect or unsafe clinical decisions. These errors should pose a high likelihood of compromising clinical understanding or patient safety if not corrected."

Error Categories

Error category	Error	Description
Hallucinations	Fabricated claim Misleading justification Detail misidentification False comparison Incorrect recommendation	Introduction of a claim not present in the input. Incorrect reasoning, leading to misleading conclusions. Incorrect reference to a detail in the input. Mentioning a comparison not supported by the input. Suggesting a diagnosis/follow-up outside the input.
Omissions	Missing claim Missing comparison Missing context	Failure to mention a claim present in the input. Omitting a comparison that details change over time. Omitting details necessary for claim interpretation.
Certainty Misalignments	Overstating intensity Understating intensity	Exaggerating urgency, severity, or confidence. Understating urgency, severity, or confidence.
Other	Other	Additional errors not covered.

MedVAL-Bench

- A dataset for **training** and **evaluation** of medical text validators
- Contains: (1) ***inputs***, (2) ***outputs***, (3) ***physician assessments (only test)***

Task Name	Dataset	Task Description	$y_{ref}?$	# Train	# Test
medication2answer	MedicationQA	medication question → answer	✓	500	135
query2question	MeQSum	patient query → health question	✓	500	120
report2impression	Open-i	findings → impression	✓	500	190
radiology2simplified	Open-i	findings → patient-friendly	✗	500	—
radiology2simplified [†]	MIMIC-IV	impression → patient-friendly	✗	—	190
bhc2spanish [†]	MIMIC-IV-BHC	hospital course → spanish	✗	—	120
dialogue2note [†]	ACI-Bench	doctor-patient dialogue → note	✓	—	85
Total				2000	840

-
1. Partially open-source
2. Out-of-distribution

-
1. Fully open-source
2. In-distribution

MedVAL-Bench - Tasks

Task	Input → output	Instructional prompt
medication2answer	medication question → answer	“Answer the following medication-related patient health question.”
query2question	patient query → health question	“Summarize the patient health query into one question of 15 words or less.”
report2impression	findings → impression	“Summarize the radiology report findings into an impression with minimal text.”
report2simplified	findings → patient-friendly	“Create a simplified, patient-friendly version of the input.”
impression2simplified	impression → patient-friendly	“Create a simplified, patient-friendly version of the input.”
bhc2spanish	hospital course → spanish	“Translate the brief hospital course into Spanish.”
dialogue2note	doctor-patient dialogue → note	“Summarize the doctor/patient dialogue into an assessment and plan.”

MedVAL-Bench - Physician Study

- Your goal is to verify the existing evaluation report and revise it if necessary.
- Only categorize a factual claim as a clinically significant error if it affects clinical understanding, decision-making, or safety.
- You may reassign, revise, add, or remove errors based on your judgment.

Impression (Expert-Written) ←

Claim 1: Low left lung volume with surrounding pleural thickening and calcified pleural plaques, consistent with prior asbestos exposure.

Claim 2: Recommend comparison with prior for change in pleural thickening

Patient-Friendly Impression (Model-Generated) ←

Claim 1: Your left lung is severely damaged, and the surrounding tissue is turning into bone due to asbestos exposure.

Claim 2: This condition is rapidly worsening and will likely require immediate surgery to remove the affected lung.

Claim 3: There's no need to compare this with previous scans because the damage is already too advanced to reverse.

Clinically Significant Errors ←

Error 1: "Your left lung is severely damaged, and the surrounding tissue is turning into bone due to asbestos exposure." - Hallucination (fabricated claim): The reference mentions pleural thickening and calcified pleural plaques but does not state that the surrounding tissue is turning into bone or that the lung is severely damaged.

Error 2: "This condition is rapidly worsening and will likely require immediate surgery to remove the affected lung." - Hallucination (incorrect recommendation): The reference does not mention rapid worsening or suggest surgery as a necessary intervention.

Error 3: "There's no need to compare this with previous scans because the damage is already too advanced to reverse." - Certainty misalignment (overstating intensity): The reference explicitly recommends comparison with prior imaging for changes in pleural thickening, and the candidate dismisses this recommendation with an exaggerated claim

MedVAL-Bench - Physician Study

Overall Quality Rating

Specify the level that best matches the candidate's factual consistency with the reference.

◆ Level 1: Fully Factually Consistent 

- No hallucinations, omissions, or certainty misalignments. All factual claims match the reference.

◆ Level 2: Low-Risk Errors 

- Subtle errors such as mild overstatements or omissions with low clinical impact. The main message is unchanged.

◆ Level 3: Moderate-Risk Errors 

- Errors span at least two categories. At least one error significantly impacts clinical interpretation.

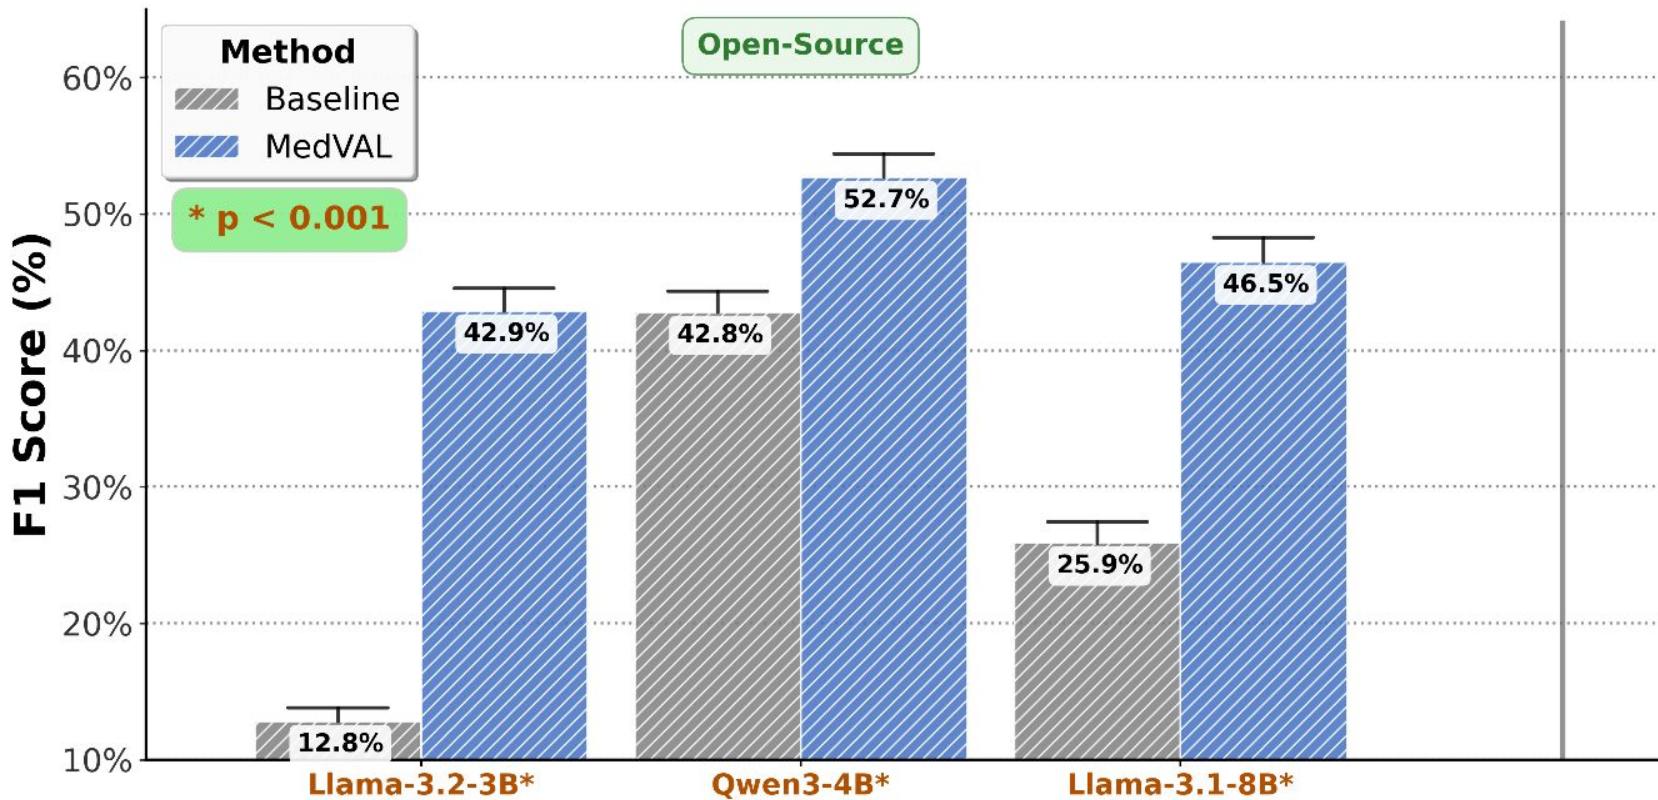
◆ Level 4: High-Risk Errors 

- Severe factual inconsistencies across all categories. At least two errors pose high clinical risk or misinterpretation.

Level 4 

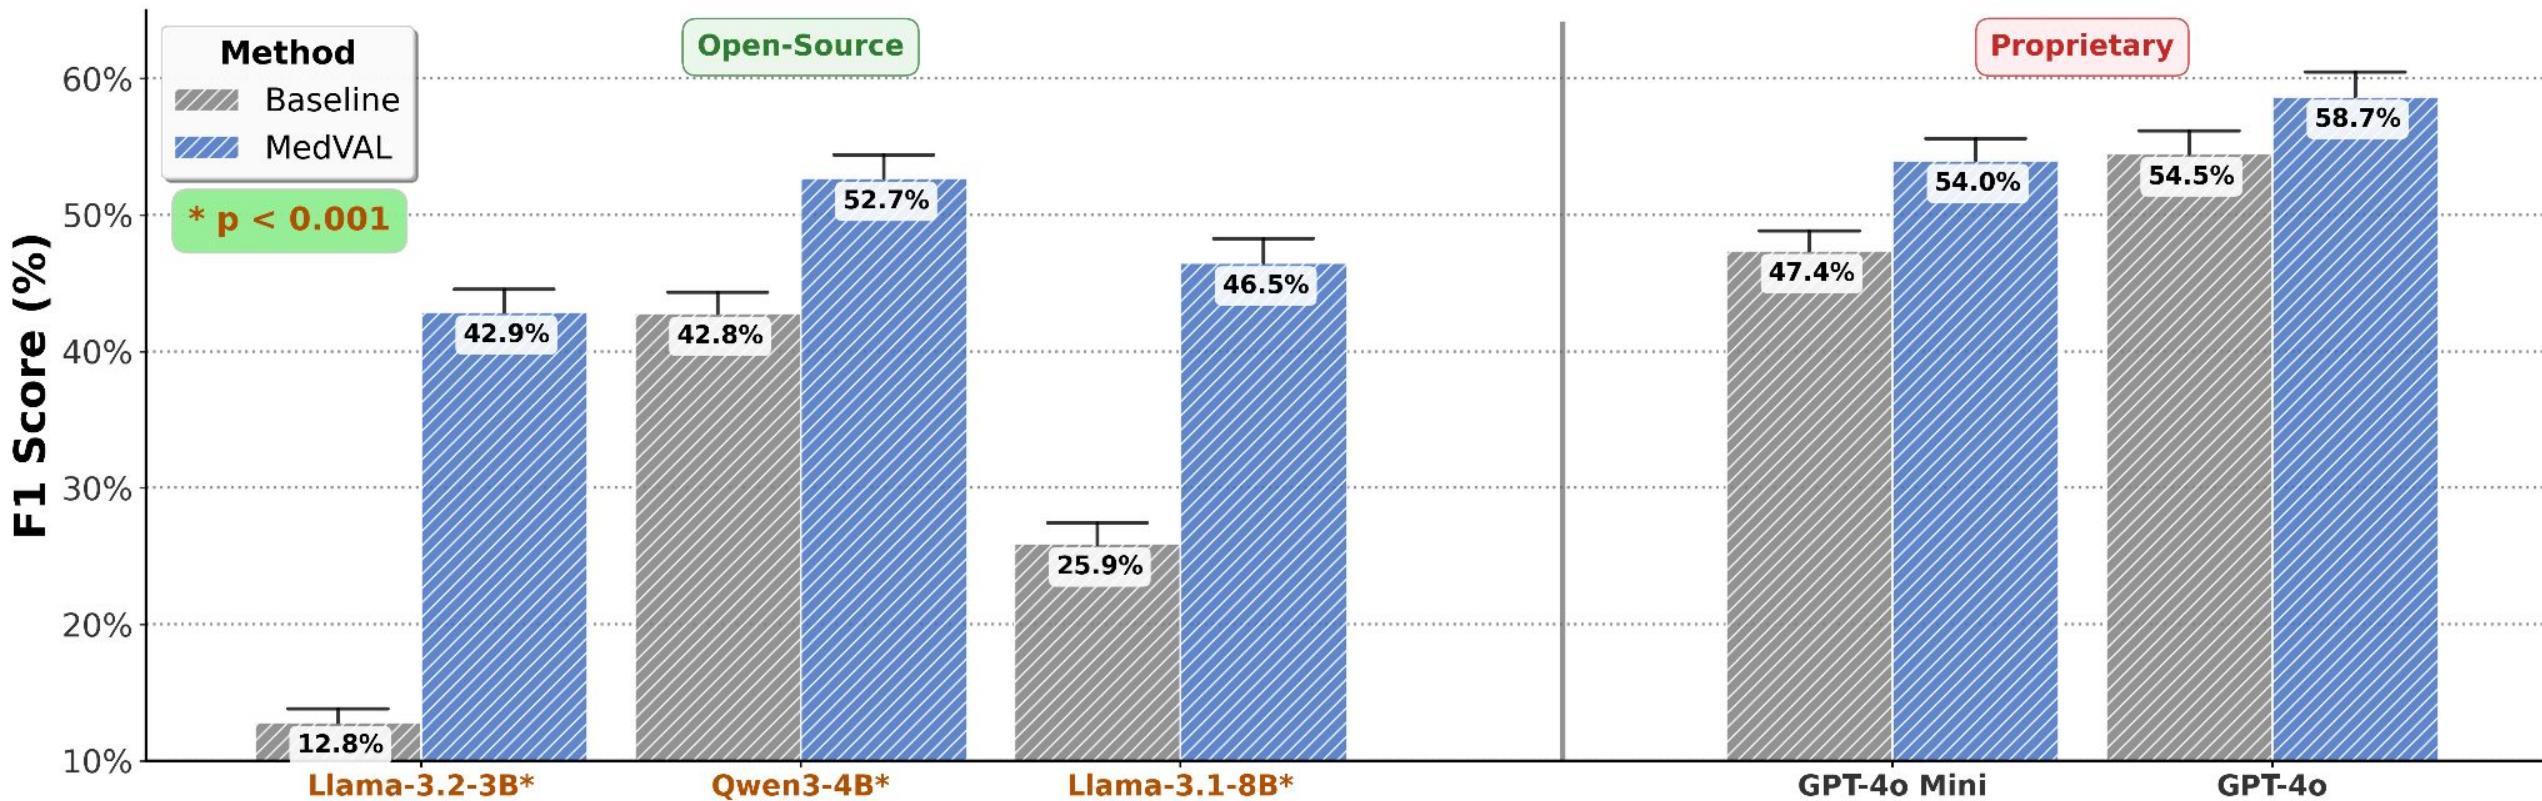
Overall Performance

a) MedVAL vs Baseline



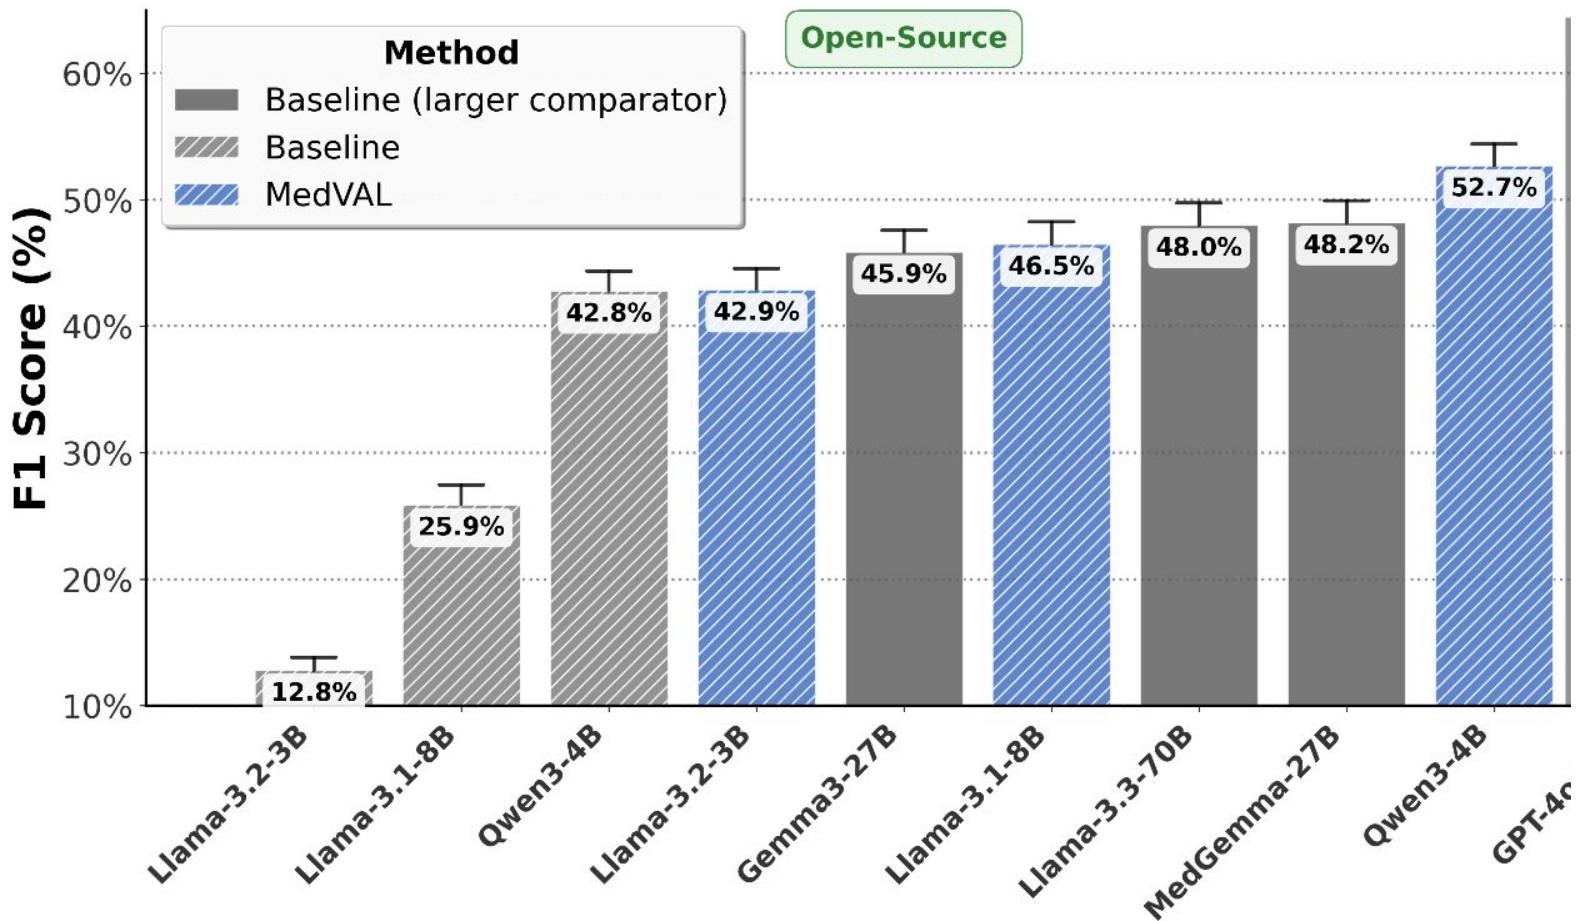
Overall Performance

a) MedVAL vs Baseline



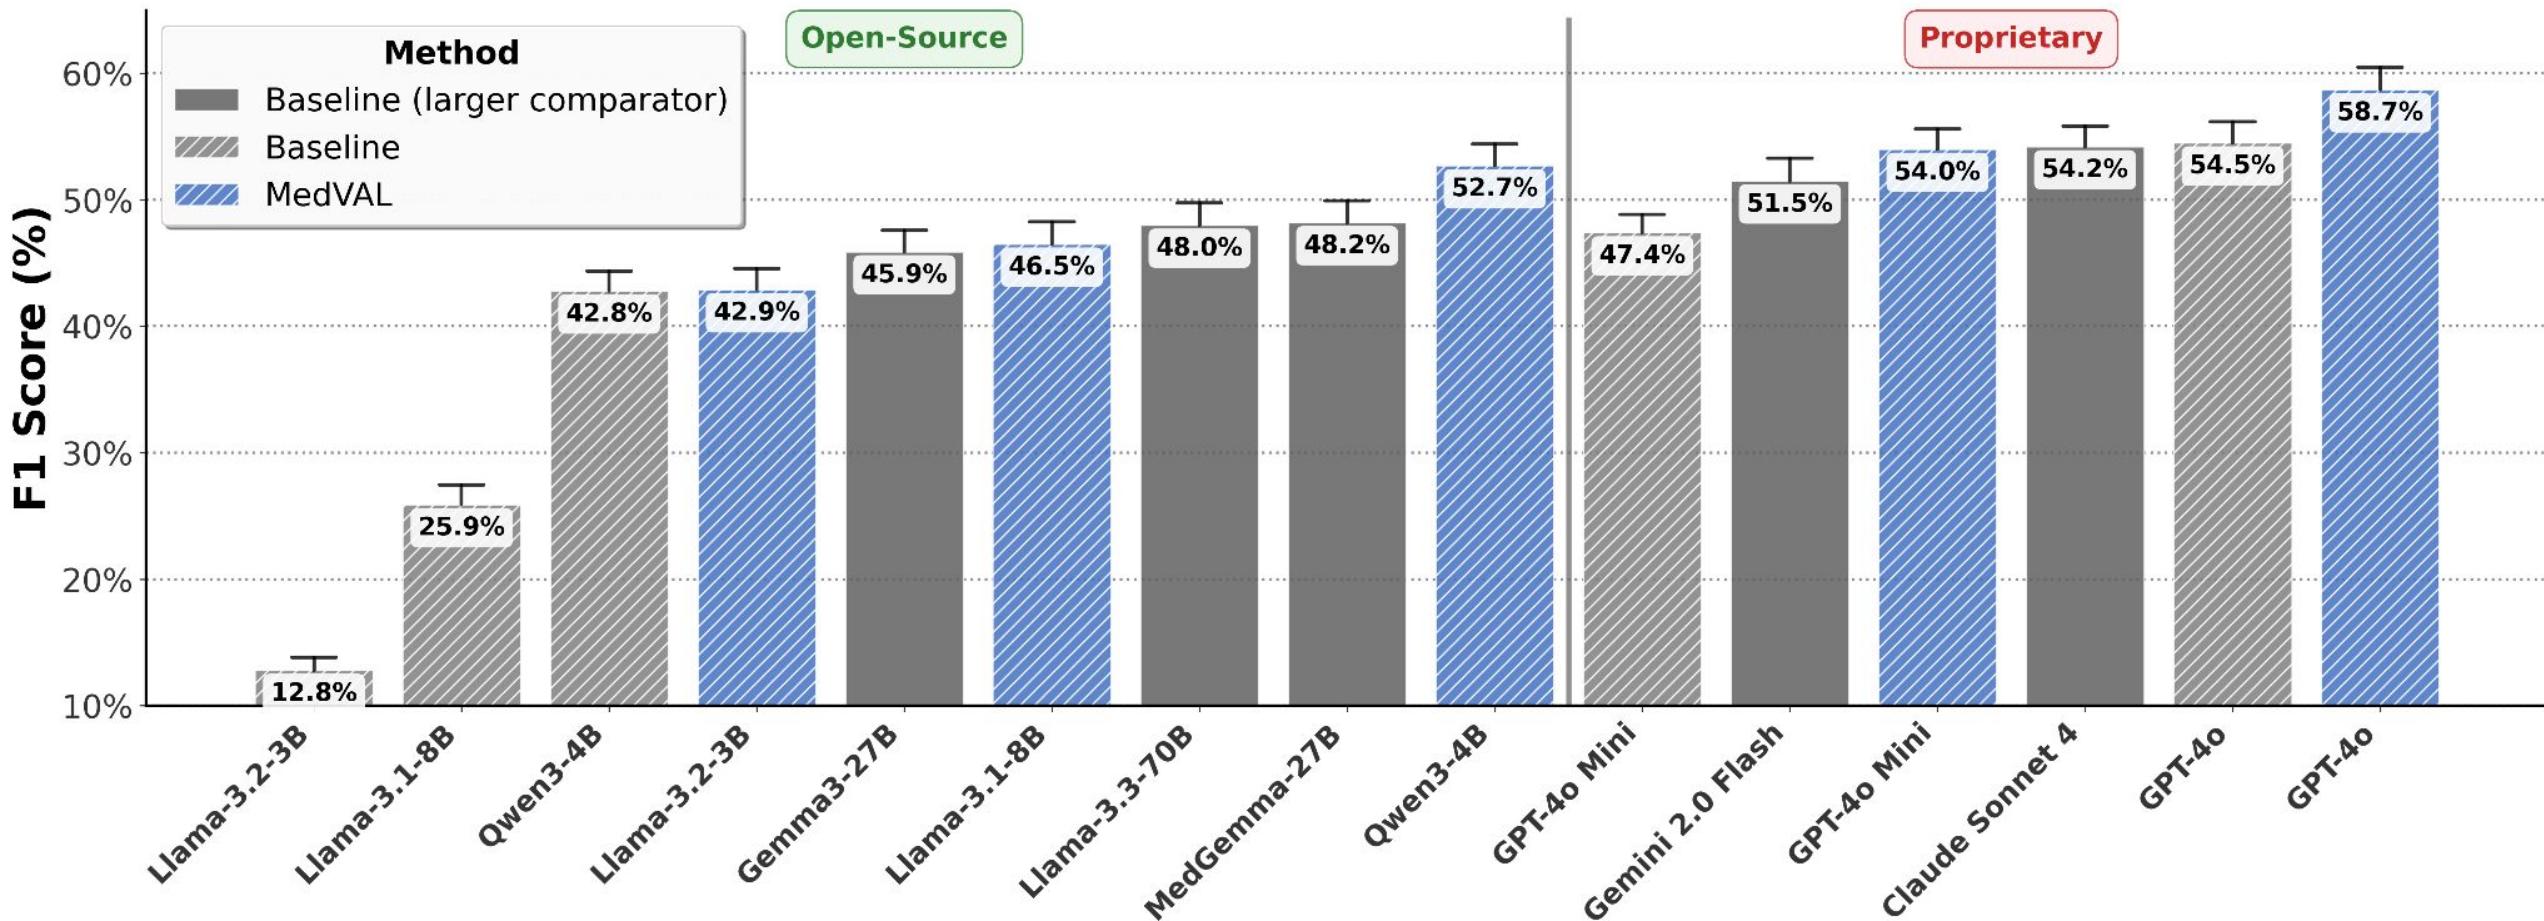
Overall Performance

b) Overall ranking



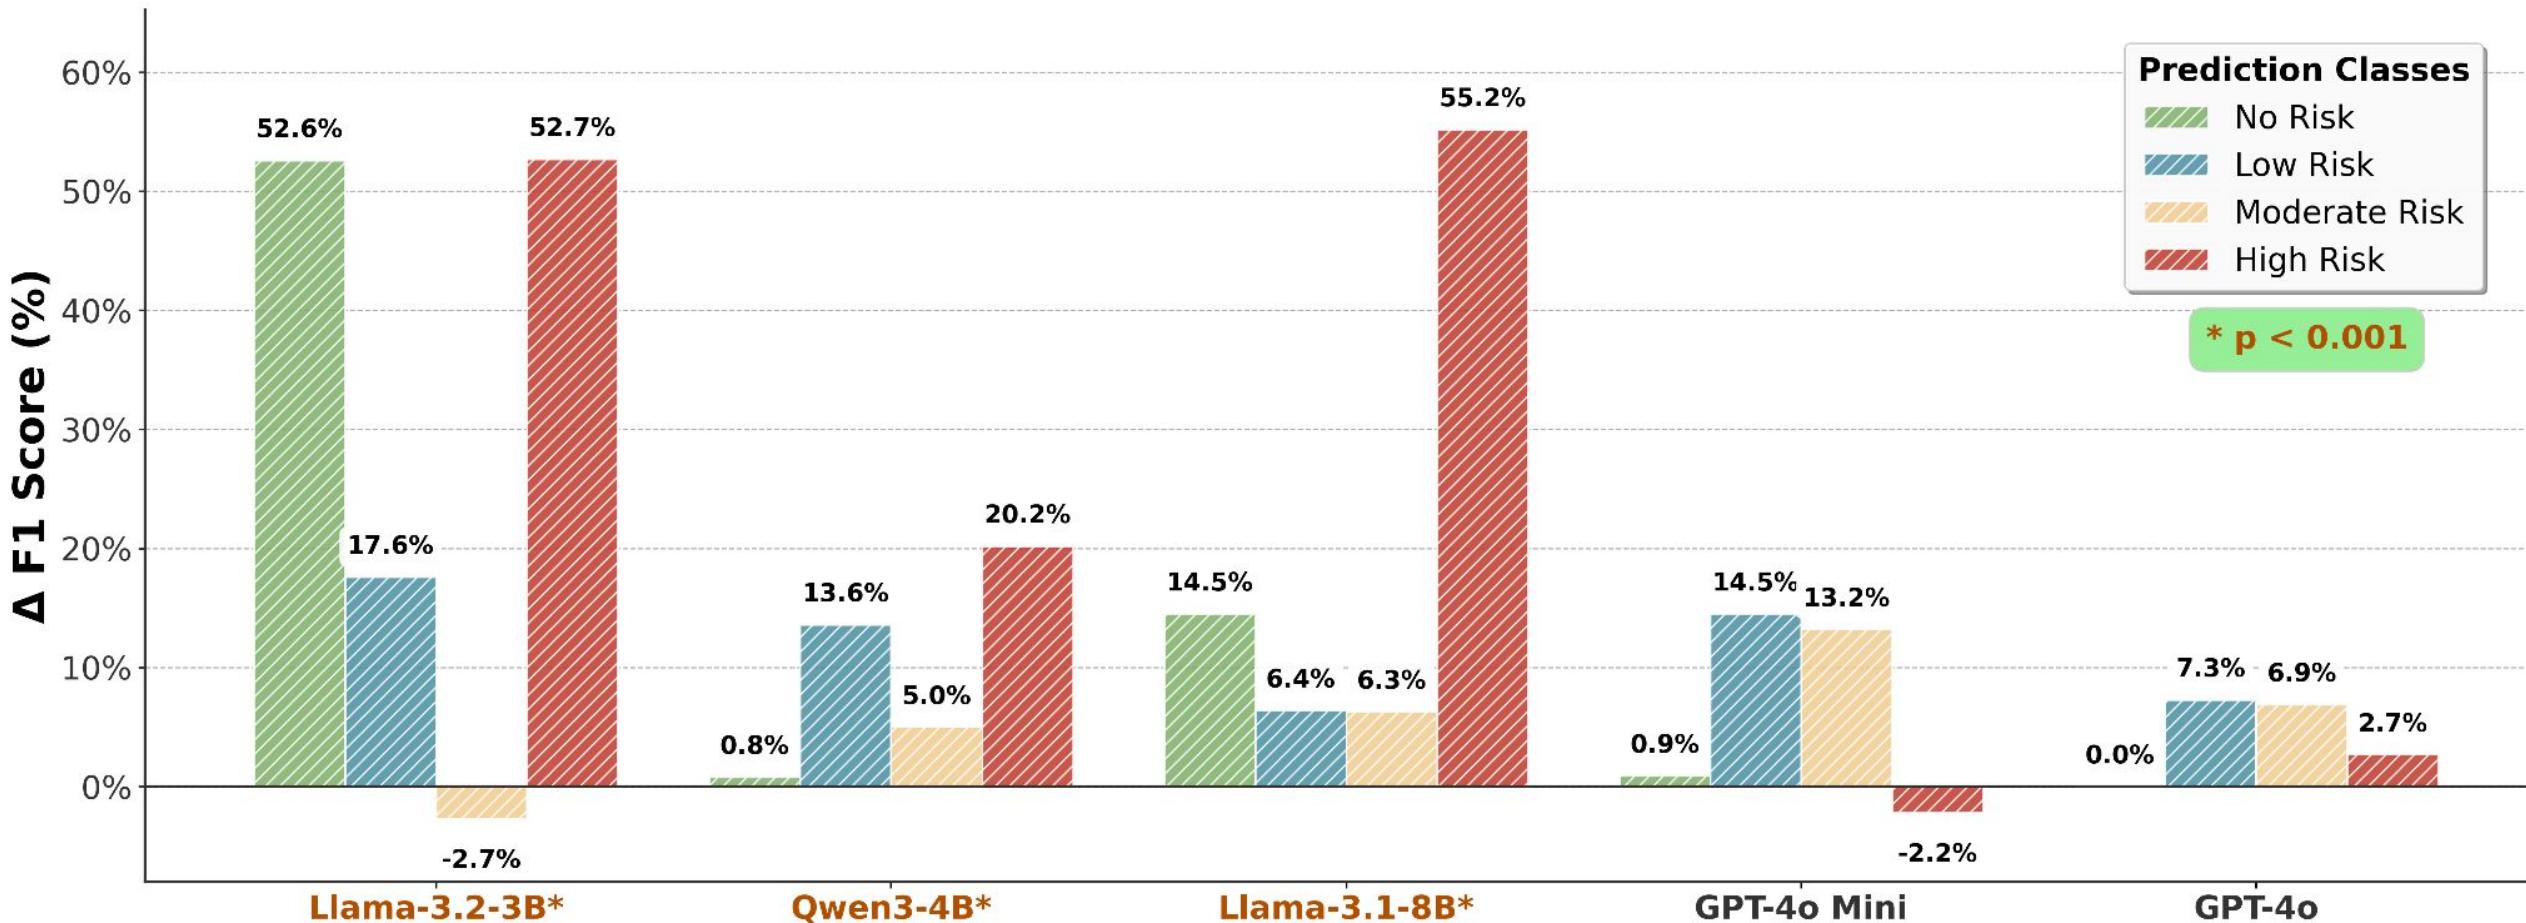
Overall Performance

b) Overall ranking



Risk-Level Classification Performance

c) Classification improvement Δ via MedVAL



Task-Wise Performance

Method	Model	In-Distribution			Out-of-Distribution			Overall
		medication2 answer	query2 question	report2 impression	impression2 simplified	bhc2 spanish	dialogue2 note	
Open-Source								
Baseline	Llama-3.2-3B	0.091	0.110	0.174	0.096	0.120	0.146	0.128
	Qwen3-4B	0.357	0.299	0.530	0.390	0.364	0.552	0.428
	Llama3.1-8B	0.342	0.285	0.278	0.225	0.158	0.113	0.259
	Gemma3-27B	0.398	0.279	0.584	0.442	0.369	0.552	0.459
	MedGemma-27B	0.462	0.287	0.616	0.451	0.349	0.603	0.482
	Llama-3.3-70B	0.478	0.311	0.633	0.496	0.362	0.322	0.480
MedVAL	Llama-3.2-3B	0.382 +320%	0.262 +138%	0.578 +232%	0.429 +347%	0.242 +102%	0.448 +207%	0.429 +235%
	Qwen3-4B	0.557 +56%	0.374 +25%	0.562 +6%	0.537 +38%	0.424 +16%	0.490 -11%	0.527 +23%
	Llama-3.1-8B	0.456 +33%	0.372 +31%	0.480 +73%	0.540 +140%	0.384 +143%	0.376 +233%	0.465 +80%

Task-Wise Performance

Method	Model	In-Distribution			Out-of-Distribution			Overall
		medication2 answer	query2 question	report2 impression	impression2 simplified	bhc2 spanish	dialogue2 note	
Open-Source								
Baseline	Llama-3.2-3B	0.091	0.110	0.174	0.096	0.120	0.146	0.128
	Qwen3-4B	0.357	0.299	0.530	0.390	0.364	0.552	0.428
	Llama3.1-8B	0.342	0.285	0.278	0.225	0.158	0.113	0.259
	Gemma3-27B	0.398	0.279	0.584	0.442	0.369	0.552	0.459
	MedGemma-27B	0.462	0.287	0.616	0.451	0.349	0.603	0.482
	Llama-3.3-70B	0.478	0.311	0.633	0.496	0.362	0.322	0.480
MedVAL	Llama-3.2-3B	0.382 +320%	0.262 +138%	0.578 +232%	0.429 +347%	0.242 +102%	0.448 +207%	0.429 +235%
	Qwen3-4B	0.557 +56%	0.374 +25%	0.562 +6%	0.537 +38%	0.424 +16%	0.490 -11%	0.527 +23% 
	Llama-3.1-8B	0.456 +33%	0.372 +31%	0.480 +73%	0.540 +140%	0.384 +143%	0.376 +233%	0.465 +80%
Proprietary								
Baseline	GPT-4o Mini	0.479	0.352	0.445	0.503	0.427	0.586	0.474
	GPT-4o	0.598	0.360	0.519	0.587	0.439	0.618	0.545
	Claude Sonnet 4	0.569	0.413	0.497	0.583	0.552	0.550	0.542
	Gemini 2.0 Flash	0.485	0.401	0.588	0.486	0.497	0.602	0.515
MedVAL	GPT-4o Mini	0.512 +7%	0.308 -13%	0.635 +43%	0.571 +14%	0.386 -10%	0.692 +18%	0.540 +14%
	GPT-4o	0.695 +16%	0.361 +0%	0.564 +9%	0.605 +3%	0.483 +10%	0.673 +9%	0.587 +8% 
Krippendorff's α								
Inter-Physician Agreement		0.904	0.560	0.861	0.872	0.943	0.830	0.848

Safety (Binary) Classification Performance

#	Model	Method	Sensitivity	Specificity	F1 Score	Accuracy
1	Llama-3.2-3B	Baseline	0.086 ± 0.01	0.960 ± 0.01	0.153 ± 0.02	0.474 ± 0.02
		MedVAL	0.919 ± 0.01	0.560 ± 0.02	0.809 ± 0.01	0.760 ± 0.01
2	Llama-3.1-8B	Baseline	0.670 ± 0.02	0.651 ± 0.03	0.688 ± 0.02	0.662 ± 0.02
		MedVAL	0.788 ± 0.02	0.786 ± 0.01	0.804 ± 0.01	0.787 ± 0.01
3	Qwen3-4B	Baseline	0.858 ± 0.02	0.643 ± 0.03	0.800 ± 0.01	0.762 ± 0.02
		MedVAL	0.839 ± 0.02	0.752 ± 0.02	0.823 ± 0.01	0.800 ± 0.01
Ensemble (1+2+3)		MedVAL	0.899 ± 0.02	0.686 ± 0.03	0.837 ± 0.01	0.805 ± 0.01 

Safety (Binary) Classification Performance

#	Model	Method	Sensitivity	Specificity	F1 Score	Accuracy
1	Llama-3.2-3B	Baseline	0.086±0.01	0.960 ±0.01	0.153±0.02	0.474±0.02
		MedVAL	0.919 ±0.01	0.560±0.02	0.809 ±0.01	0.760 ±0.01
2	Llama-3.1-8B	Baseline	0.670±0.02	0.651±0.03	0.688±0.02	0.662±0.02
		MedVAL	0.788 ±0.02	0.786 ±0.01	0.804 ±0.01	0.787 ±0.01
3	Qwen3-4B	Baseline	0.858 ±0.02	0.643±0.03	0.800±0.01	0.762±0.02
		MedVAL	0.839±0.02	0.752 ±0.02	0.823 ±0.01	0.800 ±0.01
Ensemble (1+2+3)		MedVAL	0.899±0.02	0.686±0.03	0.837 ±0.01	0.805 ±0.01
4	GPT-4o Mini	Baseline	0.784±0.02	0.807±0.02	0.809±0.02	0.794±0.02
		MedVAL	0.848 ±0.02	0.831 ±0.02	0.855 ±0.01	0.840 ±0.01
5	GPT-4o	Baseline	0.835 ±0.02	0.861±0.02	0.858 ±0.01	0.846 ±0.01
		MedVAL	0.792±0.02	0.906 ±0.02	0.849±0.01	0.843±0.01
Ensemble (4+5)		MedVAL	0.874 ±0.02	0.815±0.02	0.864 ±0.01	0.848 ±0.01

Krippendorff's α

Example: Impression (Input) -> Patient Friendly (Output)

Input (impression):

1. Interval resolution of heterogeneous ground-glass opacity on CT of the lungs. 2. Extensive persistent air trapping, has slightly resolved since prior imaging. The overall appearance is consistent with airways disease, probably asthma. No evidence of respiratory bronchiolitis - interstitial lung disease.

Color key:

Risk grade prediction:

Green: correct

Red: incorrect

Error categories:

Purple: fabricated claim

Blue: detail misidentification

Brown: misleading justification

Orange: no errors found.

Example: Impression (Input) -> Patient Friendly (Output)

Input (impression):

1. Interval resolution of heterogeneous ground-glass opacity on CT of the lungs. 2. Extensive persistent air trapping, has slightly resolved since prior imaging. The overall appearance is consistent with airways disease, probably asthma. No evidence of respiratory bronchiolitis - interstitial lung disease.

Output (patient-friendly impression):

Your recent CT scan shows that some of the cloudy areas in your lungs have cleared up. There is still some trapped air in your lungs, but it has improved a little since your last scan. This suggests you might have asthma or another breathing condition. There is no sign of a more serious lung disease like interstitial lung disease, which is good news. However, we may need to monitor for signs of respiratory infections in the future.

Color key:

Risk grade prediction:	
Green:	correct
Red:	incorrect
Error categories:	
Purple: fabricated claim	
Blue: detail misidentification	
Brown: misleading justification	
Orange: no errors found.	

Example: Impression (Input) -> Patient Friendly (Output)

Input (impression):

1. Interval resolution of heterogeneous ground-glass opacity on CT of the lungs. 2. Extensive persistent air trapping, has slightly resolved since prior imaging. The overall appearance is consistent with airways disease, probably asthma. No evidence of respiratory bronchiolitis - interstitial lung disease.

Output (patient-friendly impression):

Your recent CT scan shows that some of the cloudy areas in your lungs have cleared up. There is still some trapped air in your lungs, but it has improved a little since your last scan. This suggests you might have asthma or another breathing condition. There is no sign of a more serious lung disease like interstitial lung disease, which is good news. However, we may need to monitor for signs of respiratory infections in the future.

Assessment (physician):

Risk grade: Level 3 (moderate risk)

Errors: "However, we may need to monitor for signs of respiratory infections in the future" is missing from the input.

Color key:

Risk grade prediction:

Green: correct
Red: incorrect

Error categories:

Purple: fabricated claim
Blue: detail misidentification
Brown: misleading justification
Orange: no errors found.

Example: Impression (Input) -> Patient Friendly (Output)

Input (impression):

1. Interval resolution of heterogeneous ground-glass opacity on CT of the lungs. 2. Extensive persistent air trapping, has slightly resolved since prior imaging. The overall appearance is consistent with airways disease, probably asthma. No evidence of respiratory bronchiolitis - interstitial lung disease.

Output (patient-friendly impression):

Your recent CT scan shows that some of the cloudy areas in your lungs have cleared up. There is still some trapped air in your lungs, but it has improved a little since your last scan. This suggests you might have asthma or another breathing condition. There is no sign of a more serious lung disease like interstitial lung disease, which is good news. However, we may need to monitor for signs of respiratory infections in the future.

Assessment (physician):

Risk grade: Level 3 (moderate risk)

Errors: "However, we may need to monitor for signs of respiratory infections in the future" is missing from the input.

Assessment (baseline):

Risk grade: Level 1 (no risk)

Errors: None.

Color key:

Risk grade prediction:

Green: correct

Red: incorrect

Error categories:

Purple: fabricated claim

Blue: detail misidentification

Brown: misleading justification

Orange: no errors found.

Example: Impression (Input) -> Patient Friendly (Output)

Input (impression):

1. Interval resolution of heterogeneous ground-glass opacity on CT of the lungs. 2. Extensive persistent air trapping, has slightly resolved since prior imaging. The overall appearance is consistent with airways disease, probably asthma. No evidence of respiratory bronchiolitis - interstitial lung disease.

Output (patient-friendly impression):

Your recent CT scan shows that some of the cloudy areas in your lungs have cleared up. There is still some trapped air in your lungs, but it has improved a little since your last scan. This suggests you might have asthma or another breathing condition. There is no sign of a more serious lung disease like interstitial lung disease, which is good news. However, we may need to monitor for signs of respiratory infections in the future.

Assessment (physician):

Risk grade: Level 3 (moderate risk)

Errors: "However, we may need to monitor for signs of respiratory infections in the future" is missing from the input.

Assessment (baseline):

Risk grade: Level 1 (no risk)

Errors: None.

Assessment (MedVAL):

Risk grade: Level 3 (moderate risk)

Errors: "However, we may need to monitor for signs of respiratory infections in the future" is not mentioned in the input.

Color key:

Risk grade prediction:	
Green:	correct
Red:	incorrect
Error categories:	
Purple: fabricated claim	
Blue: detail misidentification	
Brown: misleading justification	
Orange: no errors found.	

Example: Dialogue (Input) -> Assessment and Plan (Output)

Input (doctor-patient dialogue):

[doctor] hey diana it's good to see you in here so i see that you injured your wrist could you tell me a bit about what happened [patient] yeah i was walking up and down the stairs i was doing my laundry and i slipped and i tried to catch myself ... joint dorsal angulation of the distal fracture fragment is present to a variable degree if dorsal angulation is severe presenting with a dinner fork deformity ulnar styloid fracture is present.

Output (assessment and plan):

Assessment: The patient presents with a right wrist injury sustained during a fall while doing laundry ... elevated blood pressure and heart rate are noted, possibly indicating a systemic issue.

Plan: 1. Apply a thumb spica brace for immobilization, to be worn for eight weeks ... 7. Follow up in one month to reassess progress and determine readiness for travel.

Assessment (physician):

Risk grade: Level 4 (high risk)

Errors: 1. The reference describes the fracture as extra-articular, not intra-articular. 2. The reference specifies a brace to be worn for six weeks.

Assessment (baseline):

Risk grade: Level 3 (moderate risk)

Errors: Insufficient information about patient's overall health and comorbidities.

Assessment (MedVAL):

Risk grade: Level 4 (high risk)

Errors: 1. The reference states the fracture is extra-articular, not intra-articular. 2. The reference specifies a brace (not a thumb spica) to be worn for six weeks.

Color key:

Risk grade prediction:

Green: correct

Red: incorrect

Error categories:

Purple: fabricated claim

Blue: incorrect recommendation

Brown: missing context

Orange: no errors

Conclusion

- We introduce MedVAL, a **generalizable, self-supervised** framework for **validating LM-generated medical text**
- Across all settings, MedVAL **improved average F1 scores for all underlying models**
- Risk-level analysis revealed that **MedVAL enhances model sensitivity**, particularly at **intermediate risk levels (2–3)**, which are critical for deciding human review.
- Task-wise results confirmed **strong generalization across in-distribution and out-of-distribution settings**
- Notably, MedVAL **displayed strong improvements on dialogue2note**, the **longest input context** (average 1.5k tokens) **out-of-distribution task**, demonstrating **robustness on challenging, real-world medical tasks**.

Open-Source

- Paper: <https://arxiv.org/abs/2507.03152>
- Code: <https://github.com/StanfordMIMI/MedVAL>
- MedVAL-Bench Dataset: <https://huggingface.co/datasets/stanfordmimi/MedVAL-Bench>
- MedVAL-4B Model: <https://huggingface.co/stanfordmimi/MedVAL-4B>

Open-Source

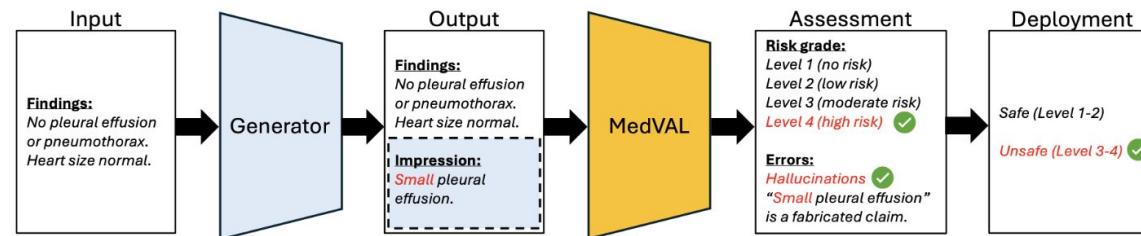
- Paper: <https://arxiv.org/abs/2507.03152>
- Code: <https://github.com/StanfordMIMI/MedVAL>
- MedVAL-Bench Dataset: <https://huggingface.co/datasets/stanfordmimi/MedVAL-Bench>
- MedVAL-4B Model: <https://huggingface.co/stanfordmimi/MedVAL-4B>

stanfordmimi/MedVAL-4B like 0 Following stanford MIMI Group 29

Text Classification Transformers Safetensors stanfordmimi/MedVAL-Bench English

Model card Files and versions Community Settings

MedVAL-4B (medical text validator) is a language model fine-tuned to **assess AI-generated medical text outputs at near physician-level reliability**.



Confidential

50

Open-Source

- Paper: <https://arxiv.org/abs/2507.03152>
- Code: <https://github.com/StanfordMIMI/MedVAL>
- MedVAL-Bench Dataset: <https://huggingface.co/datasets/stanfordmimi/MedVAL-Bench>
- MedVAL-4B Model: <https://huggingface.co/stanfordmimi/MedVAL-4B>

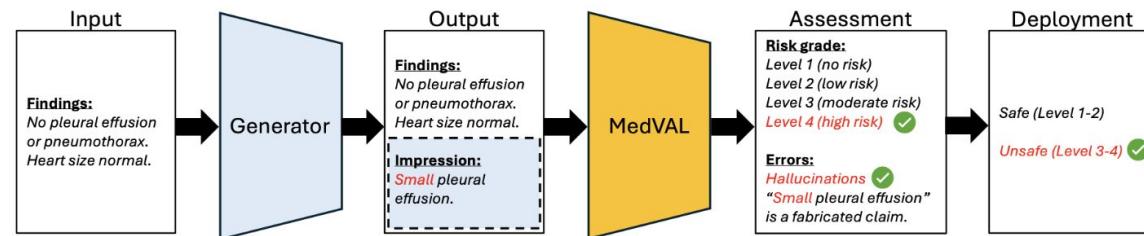
stanfordmimi/MedVAL-4B

like 0 Following Stanford MIMI Group 29

Text Classification Transformers Safetensors stanfordmimi/MedVAL-Bench English

Model card Files and versions Community Settings

MedVAL-4B (medical text validator) is a language model fine-tuned to **assess AI-generated medical text outputs at near physician-level reliability.**



Datasets: stanfordmimi/MedVAL-Bench

like 0 Following Stanford MIMI Group 29

Tasks: Text Classification Text Generation Modalities: Tabular Text Formats: csv Languages: English Spanish

Libraries: Datasets pandas Croissant +1 License: mit

Dataset card Data Studio Files and versions Community 1 Settings

Dataset Viewer

Split (2)
train · 2k rows

Search this dataset

#	int64	id	int64	task	string · classes	input	string · lengths	reference_output	string · lengths	output	string · classes	phys
1	1	2k	1	500	4 values	10	1.91k	4	4.32k ø	1 value	1 value	1 v
1	1	1	1	1	medication2answer	vitamin b strong compound is used...	na	For the treatment of clinical and sub-clinical...	na	na	na	na
2	2	2	2	2	medication2answer	what is barbiturates	na	Barbiturates are sedative-hypnotic drugs that were...	na	na	na	na

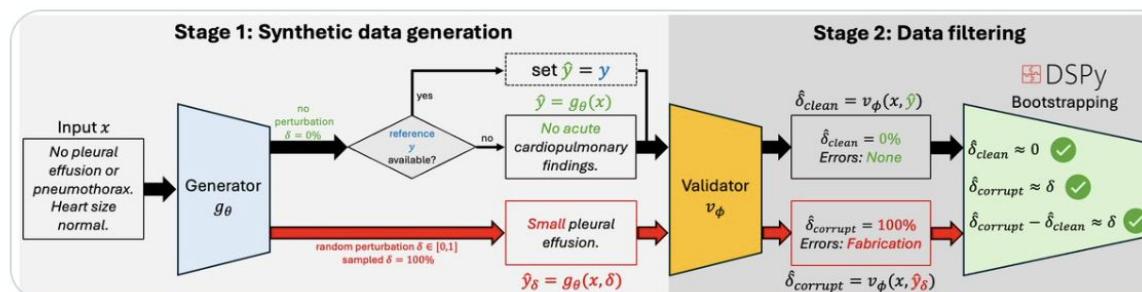
New paper from Stanford University.

"Expert-level validation of AI-generated medical text with scalable language models"

The authors use `dspy.BootstrapFinetune` for offline RL to update the weights of their LLMs.

They introduce MedVAL, a method to train LLMs to evaluate whether LM-generated medical outputs are factually consistent with inputs.

As part of their contribution, they wrote a PR to extend DSPy's parameter-efficient fine-tuning optimizers to enable Quantized Low-Rank Adaptation (QLoRA).



11:34 AM · Jul 10, 2025 · 10.8K Views



Asad Aali, MS

Research Scientist
Stanford University